



**Politecnico
di Torino**

Politecnico di Torino

Corso di Laurea Magistrale

A.a. 2025/2026

Graduation Session Mese Anno

Software Design and Architecture

Docenti:

Antonio Vetro'
Marco Torchiano

Studente:

Cristiano Berardo

Table of Contents

1	Introduction to Software Design	1
1.1	Complexity	1
1.1.1	Syntax complexity	1
1.1.2	Cause of complexity	1
1.1.3	Strategical vs. Tactical Programming	1
1.2	Modules	1
1.2.1	Abstraction	1
1.2.2	Information Hiding	2
1.2.3	Information Leakage	2
1.2.4	Information Hiding Red Flags	2
1.2.5	Method separation	2
1.2.6	Do one simple thing and complexity	2
1.3	Obscurity	2
1.3.1	Abstraction layers	2
1.3.2	Writing Comments	2
1.3.3	Naming	3
1.4	Errors and Anomalies	3
1.4.1	Define Errors out of Existence	3
1.4.2	Limit exceptions	3
1.5	Dependency Dimensions	3
1.5.1	Structural (Code-Level) Dependencies	3
1.5.2	Runtime / Distribution Dependencies	3
1.5.3	Data-Level Dependencies	3
1.5.4	Behavioral Dependencies	3
2	Design Patterns	4
2.1	Types of Software Patterns	4
2.1.1	Architectural pattern	4
2.1.2	Design patterns	4
2.1.3	Idioms	4
2.2	Pattern language	5
2.3	Idioms	5
2.3.1	Strong typing vs. Duck Typing	5
2.3.2	Property	5
2.4	Design Patterns	5
2.4.1	Creational patterns	5
2.4.2	Structural patterns	6
2.4.3	Behavioral patterns	7
2.4.4	Template Method	8
2.4.5	Strategy Pattern	8
2.4.6	Iterator pattern	8
2.4.7	Observer pattern	8
2.4.8	Visitor	8
2.5	Inheritance vs. composition	9
2.6	Abstract Syntax Tree	9
2.7	Model View Controller	9
2.7.1	MVC Principles	9
2.7.2	Execution flow	9
2.8	Summary	9

3	What is software architecture ?	10
3.1	Goals of an architecture	10
3.2	The clear architectures	10
3.2.1	Entities	10
3.3	Use cases	11
3.4	Interface adapters	11
3.5	Frameworks and drivers	11
3.6	Separation of policy from details	11
4	From design to architecture: Design Principles	12
4.1	The SOLID design principles	12
4.1.1	Single Responsibility Principle - SRP	12
4.2	Open-Closed Principle - OCP	12
4.2.1	Liskov substitution principle - LSP	13
4.2.2	Interface segregation principle - ISP	13
4.2.3	Dependency inversion principle - DIP	14
4.2.4	DIP: practical implications	14
5	Components	15
5.1	Identify components	15
5.1.1	Putting components together	16
5.2	Component principles	16
5.2.1	Component cohesion	16
5.2.2	Component coupling	17
5.2.3	The Stable Dependencies Principle - SDP	18
5.2.4	Stable Abstractions Principle - SAP	18
5.3	Strategic decisions	19
5.3.1	Conway's law	19
5.3.2	Monolithic versus Distributed architecture	19
6	Architecture visualization: C4 notation	20
6.1	Issues in sw architecture visualization	20
6.2	C4 Model	20
6.2.1	Container	20
6.2.2	Component	20
6.2.3	Code	20
6.3	Different type of diagrams	21
6.4	Level 1: Context diagram	21
6.4.1	Information reported - People	21
6.4.2	Information reported - Software system	21
6.5	Level 2: Container diagram	22
6.5.1	Elements	22
6.5.2	Interactions between containers	22
6.6	Level 3: Components diagram	22
6.6.1	Elements	23
6.7	Level 4: Code diagram	23
6.8	Summary of C4 model	23
7	Wring Components	24
7.1	Components vs. Service	24
7.2	Framework	24
7.2.1	Inversion of control	24
7.2.2	Patterns that enable inversion of control	24
7.3	Wiring components	25
7.3.1	Interacction mode	25
7.3.2	Call binding	25
7.3.3	Method/funciton calls	25
7.4	Decoupling solutions	25
7.4.1	Dependency injection	25
7.4.2	Constructor injection	26
7.4.3	Setter injection	26
7.5	Service locator	27
7.5.1	Locator vs. Dependency injection	27
7.6	Messages/Events	28
7.6.1	Roles: Producer, Consumer	28
7.6.2	Roles: Broker	28
7.6.3	Delivery Guarantees	28
7.6.4	Structure	28

7.6.5	Shared state example	29
7.6.6	Effects	29
7.6.7	Challenges	29
7.7	Shared state	29
7.7.1	Effects	29
8	Communication	30
8.0.1	One-to-one communication	30
8.0.2	One-to-many communication	30
8.1	Messaging	30
8.1.1	Message types	31
8.2	Channel	31
8.2.1	Point-to-point channels	31
8.2.2	Publish-subscribe channels	31
8.3	Supported styles	31
8.4	Direct messaging - brokerless	31
8.5	Broker-based messaging	31
8.5.1	Broker characteristics	32
8.6	Coordination	32
8.6.1	Choreography	32
8.6.2	Orchestration	32
9	Microservices	33
9.0.1	Monolith vs. Microservices	33
9.0.2	Decomposition	33
9.0.3	Guidelines from OOP	34
9.1	Business logic	34
9.1.1	Transaction script	34
9.1.2	Domain model	35
9.2	Domain Driven Design - DDD	35
9.2.1	Sub-Domain & Bounded Context	35
9.2.2	Class patterns	35
9.2.3	Domain boundaries	35
9.2.4	Aggregate pattern	35
9.2.5	Rule: reference roots only	35
9.2.6	Rule: primary keys inter aggregate	36
9.2.7	Rule: one aggregate per transaction	36
9.3	Split monolith with aggregates	36
9.3.1	Aggregate granularity	36
9.3.2	Domain Events	36
9.4	Communication	37
9.4.1	API evolution	37
9.5	Message formats	37
9.6	Remote procedure call (RPC)	37
9.6.1	Robust Synchronous APIs	37
9.6.2	Circuit Breaker pattern	37
9.6.3	Service Discovery	37
9.7	Messaging	38
9.7.1	Competing receivers	38
9.8	Shared channels	38
9.8.1	Duplicate messages	38
9.8.2	Transactional messaging	38
9.9	Data and Transactions	38
9.9.1	Database per service	38
9.9.2	Transactions in a monolith	38
9.9.3	Distributed data consistency	39
9.9.4	Saga pattern	39
9.9.5	Domain event pattern	40
9.9.6	Event sourcing	40
9.10	Queries	40
9.10.1	API Composition	40
9.10.2	CQRS	40
9.11	External API	41
9.12	Transition to Microservices	41
9.12.1	Strangling the monolith	41
9.12.2	Transition strategies	41
9.12.3	Refactor DB	41
9.12.4	God Class	41

10 Bending Spoons seminary (Evernote)	42
10.1 Evernote	42
10.1.1 Evernote's sync system	42
10.1.2 The sync service	42
11 Boundary	44
11.1 Drawing boundaries	44
11.1.1 The clan architecture	44
11.1.2 Independence and Decoupling of layers	44
11.2 Crossing boundaries	44
11.2.1 Data is significant. Data storage is a detail	45
11.3 Layers	45
11.4 Test	45
11.4.1 Design for testability	45
11.4.2 Humble objects	46
11.5 Architectural styles	46
11.5.1 Package by layer	46
11.5.2 Package by feature	46
11.5.3 Ports and Adapters	46
11.5.4 Package by component	46
11.6 Architectural styles	47
12 Monolithic Architectural Styles	48
12.1 Architectural Quanta	48
12.2 Architectural styles	48
12.3 Layered architecture - n tier	48
12.3.1 Layered architecture - variations	49
12.3.2 Adding layers	49
12.3.3 Isolation	49
12.3.4 Open layers	49
12.3.5 Open vs close layers	49
12.4 Modular monolithic style	50
12.4.1 Layered vs modular: components organization	50
12.4.2 Variations of modular monolithic style	50
12.4.3 Advantages and disadvantages of modular monolithic style	50
12.5 Pipeline architecture	50
12.5.1 Components	51
12.6 Micro-kernel - plug-in architecture style	51
12.6.1 UI: variants	52
12.6.2 Plug-in components	52
12.7 Monolithic Styles - recap	52
13 Software Quality	54
13.1 Target entities	54
13.2 Model structure	54
13.3 Software product quality	54
13.3.1 Functional suitability	54
13.3.2 Performance efficiency	54
13.3.3 Compatibility	54
13.3.4 Interaction capability (Usability)	54
13.3.5 Reliability	55
13.3.6 Security	55
13.3.7 Maintainability	55
13.3.8 Flexibility	55
13.3.9 Sub-characteristics	55
13.3.10 Usability - User error protection	55
13.3.11 Reliability - Maturity	55
13.3.12 Maintainability - Modularity	56
13.4 Quality in Use	56
13.4.1 Quality in Use - Characteristics	56
13.4.2 Maintainability - Effectiveness	56
13.4.3 Quality in Use Influences	56
13.5 Data quality	56
13.5.1 Accuracy	57

14 Software Observability	58
14.1 OpenTelemetry	58
14.1.1 Key concepts	58
14.1.2 Distributed Tracing	58
14.1.3 Span	58
14.1.4 Trace	59
14.2 Context Propagation	59
14.3 Metrics	59
14.3.1 Metric kind	59
14.4 Profiling	60
14.4.1 Why profiling?	60
14.4.2 Sampling vs. Tracing	60
14.4.3 Profile Types	60
14.4.4 Flamegraph	61
14.5 Java Flight Recorder (JFR)	61
14.5.1 Collected Events	61
14.5.2 JFT settings	61
15 Distributed Architectural Styles	62
15.0.1 Service-based Architecture	62
15.0.2 Domain services design	62
15.0.3 API access, UI and DB options	62
15.0.4 Database logical partitioning	62
15.0.5 Service-based architecture: characteristics	63
15.1 Microservices architecture	63
15.1.1 Granularity of services	63
15.1.2 Microservices architecture: characteristics	64
15.2 Event-driven architecture	64
15.2.1 Events vs messages	64
15.2.2 Further options	64
15.2.3 Event payload	65
15.2.4 Events payload: tradeoffs	65
15.2.5 Data loss prevention	65
15.2.6 Request-reply processing	66
15.3 Orchestrated EDA	66
15.3.1 Event drive architecture: usage	66
15.3.2 EDA architecture: characteristics	66
15.4 Space-based architecture	66
15.4.1 Space-based architecture: elements	66
15.4.2 Replicated cache data: collisions	66
15.4.3 Distributed caching	67
15.4.4 Replicated vs distributed caching	67
15.5 SBA architecture - Usage	68
15.5.1 When to use:	68
15.5.2 Space based architecture: characteristics	68
15.6 Orchestration-driven service-oriented architecture	68
15.6.1 Service-oriented architecture: layers	69
15.6.2 ODSA architecture: characteristics	69
15.7 Distributed Styles Architectural Characteristics: synopsis	69
15.7.1 Structural and engineering characteristics	69
15.7.2 Operational characteristics	69
16 Deployment	71
16.1 Deployment strategies	71
16.1.1 Infrastructure deployment	71
16.1.2 Application deployment strategies	71
16.1.3 Canary	71
16.2 Deployment visualization	72

Chapter 1

Introduction to Software Design

1.1 Complexity

Complexity is anything related to the structure of a software system that makes it hard to understand and modify the system.

$$C = \sum c_p t_p$$

Where C is determined by all software parts p , c_p is the complexity of part p , and t_p is the fraction of time spent on part p .

1.1.1 Syntax complexity

- **Change amplification:** change require operating in different places
- **Cognitive load:** how much information to learn to understand and perform a task. Some times is better to write longer code to reduce the complexity and so the cognitive load.
- **Unknown unknowns:** there are partes of the system that we don't know where they are or even that they exist.

1.1.2 Cause of complexity

Dependency: When we need to understand or modify a part of the code, we need to understand or modify other parts of the code. Try to reduce the number of dependencies and the complexity of the dependencies.

Obscuiy: alla the important information should be clearly visible.

1.1.3 Strategical vs. Tactical Programming

- **Tactical programming:** focus on having code that works as fast as possible. "Move fast and break things".
- **Strategical programming:** focus on long term maintainability and scalability. "Working code is not enough".

1.2 Modules

To improve and eliminate dependency and complexity, we can use modules. A module is a part of the code that has a well defined interface and a well defined implementation.

- **Interface:** what we need to know to use the module. Formal part, possible checked by the compiler. Represents a cost.
- **Implementation:** how the module works. Ideall much more complex than the interface. Represents a benefit.

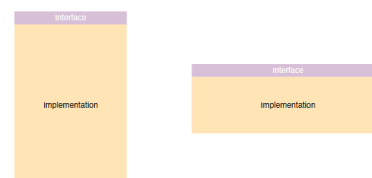
1.2.1 Abstraction

What we typically do is to perform an abstraction. Abstraction means selective removal of information. Is a simplified view of the something that omits "unimportant" information. Omitting information can lead to eads to unknowns and keeping unimportant information can leads to cognitive overload.

Deep modules and shallow interfaces

Deep modules: modules with a simple interface and a complex implementation. They are good because they reduce the cognitive load and the change amplification. They are bad because they can lead to unknown unknowns.

Shallow interfaces: modules with a complex interface and a simple implementation. They are good because they reduce the unknown unknowns. They are bad because they can lead to cognitive overload and change amplification.



1.2.2 Information Hiding

Information hiding is the principle of hiding the implementation details of a module from the users of the module. Example sort function: we want to sort a list of numbers, we don't care about how the sorting algorithm used.

1.2.4 Information Hiding Red Flags

- **temporal decomposition:** execution order is reflected in the code structure. Operations that happen in different times are in different classes.
- **overexposure:** API forces knowledge about rarely used features to use common ones. E.g. `BufferedReader r = new BufferedReader(...)`
- **Too many classes:** code is divide into too many shallow classes.

1.2.5 Method separation

Better together or Better apart?

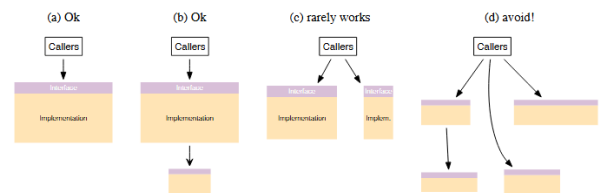
- **Bring together:** if information is shared, if we simplify the interface, if we eliminate duplication.
- **Separate:** if the code is general purpose or specific purpose.

Method separation red flags

If we have a lot of repeted code and/or we mixed specific and general code, means we have put together code that should be separated.

1.2.6 Do one simple thing and complexity

Each method should do one thing and do it completely



1.3 Obscurity

Obscurity is the principle of making the important information clearly visible and the unimportant information hidden. Obscurity is a great source of complexity. Obscure code is hard to Understand and creates more Bugs.

The solution to obscurity is writing obvious code, this means with a quick glance we can understand what the code does. We can achieve this by using abstraction layers.

1.3.1 Abstraction layers

Software systems are composed in layers. Higher layers use the facilities provided by lower layers.

Abstraction simplifies understanding and reduces obscurity.

In a well-designed system, each layer provides a different abstraction from the layers above and below it.

Abstraction Layers Red Flags

Some red flags that indicate that we have a problem with abstraction layers are:

- **Pass-Through methods:** is a method that does nothing but call another method.
- **Pass-Through variables:** is a variable that is passed through over and over down a long chain of method calls.

1.3.2 Writing Comments

Another method to reduce obscurity is writing comments. If users need to read a method code to use it, there is no abstraction.

Comments should describe things that are not obvious in the code:

- **Pick conventions:** write comments like other persons in the project
- **Don't repeat code:** avoid redundant comments that simply restate what the code already makes clear
- **Higher level comments enhance intuition**
- **Delayed comments are bad comments**
- **Use comments as design tool**

1.3.3 Naming

Togther with comments, naming is a great tool to reduce obscurity. Bad names cause bugs. Name should create a image in mind of the reader about the nature of the thing being named. Names should be precise and consistent.

1.4 Errors and Anomalies

Exceptions add complexity because they create dependencies between the code and if they are easy to throw they are hard to handle.

A piece of code might throw exceptions in different ways:

- Caller may provide bad arguments or configuration
- Invoked method may not able to complete the request
- In distributed systems network packets may be lost or delay or peers may communicate in unexpected ways
- The code may detect bugs or internal inconsistencies or situations that are not prepared to handle

1.4.1 Define Errors out of Existence

What are the solutions to errors? Define errors out of existence. This means, not ignore the error but to redefine the behavior of a method so that no error is required.

E.g. `String.substring(int beginIndex, int endIndex)` it throws different exceptions Some in only specific cases. This is completely avoidable and unnecessary.

1.4.2 Limit exceptions

There are some methods to limit the number of exceptions that a method can throw:

- Handle exceptions at the point where they are thrown, if possible
- Exception aggregation means handle different exceptions with same handler
- Just crash
- Mask exceptions avoiding unexpected exceptions

1.5 Dependency Dimensions

1.5.1 Structural (Code-Level) Dependencies

- **Implementation Dependency:** some piece of code depends on a concrete class instead of an abstraction
- **Construction Dependency:** you depend on how a collaboration, another class, is created.
- **Compile-Time Dependency:** the module knows something at compile time

1.5.2 Runtime / Distribution Dependencies

Often linked to distributed architectures.

- **Spatial Dependency:** A component depends on the location of another component. E.g. hard-coded URLs
- **Temporal Dependency:** A component depends on another being available at the same time. E.g. remote method calls or servers
- **Synchronization Dependency:** A component depends on another to complete work before continuing. E.g. synchronous request/response

1.5.3 Data-Level Dependencies

- **Data Dependency:** One component depends on another's internal data model. E.g. shared database
- **Schema Dependency:** Multiple services depend on the same database schema or representation structure. E.g. DB schema, XML/Json schema

1.5.4 Behavioral Dependencies

- **Control Flow Dependency:** One component dictates another's execution path. E.g. Centralized process managers
- **Transactional Dependency:** E.g. Components rely on atomic distributed consistency.
- **Timing Assumption Dependency:** Assumptions about response time or ordering. E.g. "Service must respond in <50ms" or Event ordering assumptions

Chapter 2

Design Patterns

Initially proposed by Christopher Alexander in the context of architecture.

A reusable **solution**
to a known **problem**
in a well defined **context**

The context is a *design* situation giving rise to a (design) problem.

The problem is a set of forces repeatedly arising in the context. ¹

The solution is a proven resolution of the problem. Configuration to balance forces. Structure with components and relationships and run-time behaviour.

Example of Social Pattern: You are in queue to the supermarket gastronomic section. The problem is hard to spot who arrived first. The solution is to use a ticket machine.

2.1 Types of Software Patterns

- **Architectural patterns** - high level structures of software systems. E.g., layered architecture, client-server, microservices.
- **Design patterns** - solutions to common design problems. E.g., Singleton, Observer, Factory Method.
- **Idioms** - language-specific patterns. E.g., RAII in C++, list comprehensions in Python.

2.1.1 Architectural pattern

Normally the architectural pattern works with big components and their interactions and is needed to describe the responsibilities of these components. Expresses a fundamental structural organization schema for software systems and defines the rules and guidelines for organizing the relationships between the components.

Example of architectural pattern: Several programs that are used in sequence read from input and write sequentially to output.

The problem is there are a lot of intermediate files used for communication between programs.

The solution is to adopt pipes and filters architecture feeding with the result of the previous one.

2.1.2 Design patterns

Provides a scheme for refining components of a software system or their relationships and describes a commonly recurring structure of communicating components.

Example of design pattern: A class library providing few functionalities contains a lot of classes.

The problem is the user is exposed to the internal complexity of the library.

The solution is to create a new façade class that interacts with the user and hide all the details.

2.1.3 Idioms

It is a low-level pattern specific to a programming language. Describes how to implement particular aspects of components or the relationships between them. Leverages the features of a programming language.

¹Any relevant aspect of the problem E.g., requirements, constraints, desirable properties

Example of idioms: An attribute is constant and should be globally available to many classes.

The problem is opening access would allow unauthorized modifications and the attribute is repeated in every object.

The solution is to make public static final in Java language.

2.2 Pattern language

Pattern do not exist in isolation in fact can happen that two or more patterns are applied together, used to implement part of another pattern and introduce a problem solved by another

They are called Pattern Languages or pattern systems.

Collection of patterns together with guidelines for:

- Implementation
- Combination
- Practical use

Should:

- Count enough patterns
- Describe patterns uniformly
- Present relationships

2.3 Idioms

X-ability

Clonability, serializability, comparability, etc. some specific behaviour dealing with many heterogeneous classes. The solution is to derive from a class or interface that defines the required methods.

Example: In Java all methods extends the root class *Object* and provides the methods *toString()*, *equals()*, *hashCode()*, etc. that can be overridden to provide specific behaviour.

2.3.1 Strong typing vs. Duck Typing

"If it looks like a duck, swims like a duck, and quacks like a duck, then it probably is a duck."

Strong typed language: You must define interfaces and classes to be able to call a specific method. Compile time type checking.

Dynamic/duck typed language: You just add a method to a class and you can call it. No compile time type checking.

2.3.2 Property

A property, more general than an attribute, is a named characteristic of an object typically accessed through well-defined operations called accessors, getters, and setters forming the public interface of the object and it is not necessarily an attribute of the object (e.g. the age can be derived from the date of birth).

2.4 Design Patterns

Description of communicating objects and classes that are customized to solve a general design problem in a particular context.

A design pattern names, abstracts, and identifies the key aspects of a common design structure that make it useful for creating a reusable object- oriented design.

		Purpose		
		Creational	Structural	Behavioral
Scope	Class	1	1	2
	Object	4	6	10

Pattern selection How to use patterns? First of all you need to know the problem you are trying to solve, knowing the patterns and then select the one that best fits your needs.

2.4.1 Creational patterns


- **Class:** Factory Method
- **Object:** Abstract Factory, Builder, Prototype, Singleton

Abstract Factory

A family of related classes can have different implementation details.

The client should not know anything about which variant they are using / creating.

Example: GUI where you have different components (buttons, text fields, etc.) and the same GUI can be applied to different platforms such as Windows, MacOS, Linux, etc.

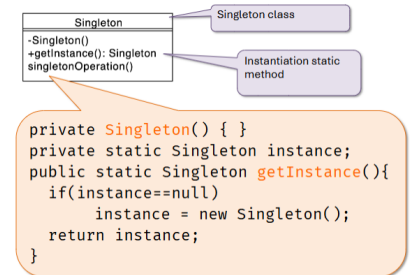


Singleton

A class represents a concept that requires a single instance.

The client should be able to access, in appropriate way, the unique instance.


Note: If we talk about of an object designer pattern, in python you can just create a global variable and it will be a singleton.



Builder

An object of a complex class has to be created.

The creation entails complex interaction with the object and could be different variations of the object.



Example: $10.40 \text{ kg} \cdot \text{m}^2 \cdot \text{s}^{-3}$

```
//Usual non-fluent builder
Measure power = new Measure(10.4);
power.addUnit("kg", 1);
power.addUnit("m", 2);
power.addUnit("s", -3);
power.setPrecision(2);
```


```
//Fluent builder
Measure power = new Measure.Builder(10.4)
    .is("kg").by("m")
    .squared().by("s").to(-3)
    .withPrecision(2).build();
```

Prototype

Several object can be created out of a palette of similar classes.

The class hierarchy can become quite complex and classes in hierarchy are differentiated by minor parameters (that could be defined through constructor arguments).

A solution could start with a prototype classe and clone it to create new objects.



Example: A graphical editor where you have very similar shapes (e.g. square with different radius, colors, etc.).

2.4.2 Structural patterns

How organize classes and objects to compose larger structures.


- **Class:** Adapter
- **Object:** Bridge, Composite, Decorator, Facade, Flyweight, Proxy

Adapter

A class provides the features required by another class but its interface is not the one expected by the client.

The integration of the provider class should be possible without modifying it. The source code could be not available or is already used as it is somewhere else.

Example: In Java, keyboard events are handled by KeyListener. Listener class need to implement the interface and implement all the methods. If you are interested only in one of the methods, you can use the KeyAdapter class that is empty and you can decide which method to override, and the Listener extends KeyAdapter instead of KeyListener.




Composite

You need to represent part-whole hierarchies of objects.

Clients can be very complex and they require to adapt the difference between all the elements in the hierarchy.


Example: In Java UI components are organized in a tree structure. You can have a panel that contains other panels and buttons, etc.



Decorator

You have some objects of classes that have specific behaviour and you want to add the same behaviour plus something else dynamically.

The solution can be extend the class but if that behaviour is something you would to apply to many classes, you need to extend all of leaf classes. Using a decorator you can modify the behaviour without creating an explosion of subclasses for every combination of features




Example: In Java we have JScrollPane can be used to provide a scroll pane to any component

Flyweight

You need need to create a very large number of similar objects, and memory usage and performance become concerns.

Example: In Java we can pool different objects together. The Integer class in Java.



Proxy


You need to controll access to an object.

We need to provide a placeholder that is able to mediate between the client needs and all the requirements you want to controll before accessing the final object without changing interface or class.

- **Remote proxy:** provides a local representative for an object in a different address space.
- **Virtual proxy:** creates expensive objects on demand. E.g. slow loading resource - images in a web page.
- **Protection proxy:** controls access to the original object. E.g. when objects should have different access rights.
- **Smart reference:** is a replacement for a bare pointer that performs additional actions when an object is accessed. E.g. counting the number of references, checking permissions, etc.

Facade

Hide the details and the complexity of a general class by providing a simpler interface. The client should not know anything about the subsystem.



2.4.3 Behavioral patterns

Behavioral patterns allow to implement specific algorithms using objects and classes to provide control and flexibility.


- **Class:** Interpreter, Template Method
- **Object:** Chain of Responsibility, Command, Iterator, Mediator, Memento, Observer, State, Strategy, Visitor

The basic mechanism of behavioral patterns is enapsultion variation. Instead of having a lot of if-else you have a structure at object level that allows you to define different alternatives.

2.4.4 Template Method

An algorithm/behavior has a stable core and several variation at given points. You have to implement/maintain several almost identical pieces of code.

Example: The compare in Java. You have to implement the compare method that solves integer or string comparison.




2.4.5 Strategy Pattern

Consequences:

Many classes or algorithm has a stable core and several behavioral variations, several different implementations are needed. Every variation of an algorithm can be embedded in an object and passed as a parameter to the algorithm.

Example: Comparator in Java.


- + Avoid conditional statements
- + Algorithms may be organized in families
- + Choice of implementations
- + Run-time binding
- Clients must be aware of different strategies
- Communication overhead
- Increased number of objects



2.4.6 Iterator pattern

You have a collection of objects and you want to iterate on it allowing multiple concurrent iterators.

Example: In Java there is the Iterable interface and Iterator.




2.4.7 Observer pattern

The change in one object may influence one or more other objects. This can cause high coupling between objects and the number and type of objects to be notified may not be known in advance.


Example: In Java there is the Observable interface and Observer interface. This is very common in UI programming. Consequences:

- + Abstract coupling between Subject and Observer
- + Support for broadcast communication
- Unanticipated updates



2.4.8 Visitor

You have an object structure that contains many classes with different interfaces and you want to perform several different and unrelated operations on these objects. The operations depend on the concrete classes and you want to avoid putting many different methods in each of these classes.




2.5 Inheritance vs. composition

Reuse can be achieved via:

- **Inheritance:** a new class is created by inheriting from an existing class and adding new features or modifying existing ones. Clients can invoke directly inherited methods.
- **Composition:** The class has a link or instance to a class you want to reuse and just forwards the calls to it.

Note: The Observable pattern is deprecated. The PropertyChangeSupport is more flexible.



2.6 Abstract Syntax Tree


The AST is a tree that represent a syntax of formal language. We can use the composite pattern to representation the arithmetic expressions.

And using the visitor pattern or the interpreter pattern to evaluate the expression.

2.7 Model View Controller

Context: Interactive applications with flexible HCI.

Problem: The same information is presented in different ways/windows, Windows must present consistent data, Should be possible to add/change presentations of information



2.7.1 MVC Principles

When building a GUI we must consider two main aspects:

- Layout (**View**): how to place the graphical elements to achieve a give visual aspect
- Events (**Controller**): which behavior associate to elements' events
- Application logic (**Model**): must remain, as far as possible, separate from user interface.





Figure 2.1: MVC

Figure 2.2: MVP



2.7.2 Execution flow

There is no predefined order of execution. Operation are performed in response to external events (e.g. mouse click). Event handling is serialized. To execute operations in parallel threads must be used. Method main in GUIs has the only goal of instantiating the graphical elements

2.8 Summary

Architectural patterns deal with overall system structure. They provide a unique metaphor for the system (e.g. pipe and filters). They address specific domains (e.g. distribution or interaction) and system evolvability

Chapter 3

What is software architecture ?

The architecture of a software system is the shape given to that system by those who build it. The form of that shape is in the division of that system into components, the arrangement of those components, and the ways in which those components communicate with each other. The purpose of that shape is to facilitate the development, deployment, operation, and maintenance of the software system contained within it.

3.1 Goals of an architecture

The software must change to adapt to software regulation, user feedback.

We have some pattern that make methods and classes easy to be changed, but now we want something behind the methods.

The goals of an architecture are :

- keep the system valuable for its users over time
- minimize the lifetime cost of the system
- maximize programmer productivity

This can be achieved by making the system :

- easy to understand
- easy to develop
- easy to maintain
- easy to deploy
- easy to change

In Java we have the package, that is a module, and is a logical grouping of classes.

Design Patterns: helps to build methods and classes inside the module in a way that they are easy to be changed.


Design Principles: define the modules can collaborate together and how they are shaped.

3.2 The clear architectures

The **Outer circles:** mechanisms, they change often.

The **Inner circles:** policies, subject to less frequent changes.

Source code dependencies point only inward, toward higher-level policies (Dependency Rule)



3.2.1 Entities

An Entity is an object within a software system that contains a set of critical business rules operating on Critical Business Data.

The interface of the Entity consists of the functions that implement the Critical Business Rules to operate on the Critical Business Data. They contain the most general and highlevel rules and neither they are often changed nor they are usually affected by lower-level changes.

Example: The system is the university portal and an important entity is the student, that has some critical business rules like "the student can enroll to the master degree only if they have completed the required prerequisites and then assign new ID".

Is the most stable part of the system and we want to preserve as much as possible from changes occurring to the other elements.

Student
- ID - DateBirth - Address
+ assignGrade() + assignSubsidee() + computeTax()

3.3 Use cases

To apply the Business Roles to the entities we need to use the use cases.

A use case is a description of the software behavior and interactions (with people or other software) in a specific case. It specifies:

- the input to be provided by the user (or other systems)
- the output to be returned
- The processing steps involved.

to determine the flow of data to and from the entities.


Software at this level change whenever a change to any element of a use case is required

Use case in practice

It represents a **functional requirement**.

- It makes clear what you expect from a system ("what?")
- It hides the behaviour of the system ("how?")
- It is a sequence of actions (with variants) that produce a result observable by an actor

The system is seen as a black-box. Use cases express the interaction (exchange of messages) between the entities that interact with the system itself (called actors).



Example - Specify annual tax for a student Input: student ID

Output: annual tax amount

Main scenario:

1. Validate ID provided by user
2. Retrieve economic level
3. Compute due past taxes
4. Compute new taxes
5. Get average grades for the last academic year
6. If average grade > 26, apply 20% discount
7. Ask operator for confirmation: if yes change student else reset

3.4 Interface adapters

Their goal is to shape data in a way that it is convenient to higher levels or other systems or artifacts interacting with:

The same is done in the opposite direction. Examples of code in this layer: Presenters for GUI, Controller for DB

Code inward to this level know anything about details (e.g., SQL)

3.5 Frameworks and drivers

In order to achieve the separation we have seen so far, we need to separate low level details from the high-level details. Layers allow independent development, deployment, operation, and maintenance.

The details are the frameworks, the database, the drivers, the web server, etc., and usually we do not write much code here.

3.6 Separation of policy from details


Decoupling of policy from the details and defer decisions on details for as long as possible (MongoDB or SQL, etc.)

DB component knows about use case rules component, but NOT viceversa.

Low level components know everything of the higher-level components, but NOT viceversa.

The Database component contains the code that translates the calls made by the BusinessRules into the query language of the database. This translation code knows about the BusinessRules.

Business rules can be adapted to any kind of DB.



Chapter 4

From design to architecture: Design Principles

Design principles means follow a clean architecture and a clean architecture means that the architecture can easily evolve, to be changed.

4.1 The SOLID design principles

The SOLID principles suggest software architects on how to arrange functions and data into groupings (e.g., classes), and how these groupings are arranged.

The goal of the SOLID principles is to create mid-level software structures (modules) that:


- Are easy to understand
- Allow changes in the most "comfortable" way
- Are the basis of reusable components

4.1.1 Single Responsibility Principle - SRP

A module should have one, and only one, reason to change. It means that a module should be responsible to **one**, and only one, **actor**.


This is helpful to separate code used by different actors and for a better maintenance.

Example *CalculateTax()* is specified by the Finance Department. *InsertGrade()* is specified by Students Office and used by Professor. *VerifyAdmissionCriteria()* is specified and used by Students Office. *CalculateTax()* and *VerifyAdmissionCriteria()* share a common algorithm for calculating the average grade.




Changes in *ComputeAverage()* affect two actors: errors may arise and also changes to Class Student due to different needs of actors may lead to collisions and merges.

Solution 1: Separate data from functions




Solution 2: Use the Façade pattern



You increase coupling and complexity, its a trade off, but the separation is splitted.

Solution 2: variation



Note: You increase coupling and complexity but the separation is splitted.

4.2 Open-Closed Principle - OCP

A software artifact should be **open for extension but closed for modification**. Using design patterns to make new functionality in a way that is additive and not to modify the existing structure.

The behavior of a software artifact ought to be extendible, without having to modify that artifact.


The aim of OCP is to make the system easy to extend without requiring too many changes to the system.

Using the open-closed principle, the software architect divides the system into components and arranges these components in a dependency hierarchy that protects higher-level components from changes in lower-level components.

Example A system displays summary information about a student's career at the university on a dashboard: Subjects (those passed with grades), Taxes paid and due, Current semester schedule, Next exams, Important notices, Stakeholders ask for a printed report with only information on subjects and taxes, and a different formatting.


Generating the report involves two separate responsibilities:

1. the calculation of the reported data
2. the presentation of that data into a web- and printer-friendly form.



4.2.1 Liskov substitution principle - LSP

If for each object $o1$ of type S there is an object $o2$ of type T such that for all programs P defined in terms of T , the behavior of P is unchanged when $o1$ is substituted for $o2$ then **S is a subtype of T** .




If we substitute $o2$ for $o1$ and the behavior of p does not change means S is derived class of T .


This principle is about subtyping, not only for inheritance or extension but also for the logic, the behavior of the class.

Derived type must fully support the substitution of their base types: functions that use pointers or references to base classes must be able to use objects of the derived without knowing it. This is related to the substitution property.

A subtype can replace its supertype only if it behaves the same way. This idea, called **behavioral subtyping**, means the subtype doesn't just have the same methods as the supertype — it also follows the same rules for how those methods should work. That way, anything the client expects from the supertype will also hold true for the subtype.



```
Rectangle r = ...;
r.setW(5); r.setH(2);
assert(r.area() == 10);
```




Depending on how r is defined, the assertion can fail because you not use the $setH()$ and $setW()$ methods of the rectangle but you use the $setSide()$ method of the square.

How to refactor it? Using the Interface segregation principle - ISP

4.2.2 Interface segregation principle - ISP

Clients should not be forced to depend on operations and/or interfaces they do not use. Better to make fine grained interfaces that are client specific. Similar to SRP we have different actors, represented by different classes. They are only interested in a subset of the operations, so we can split the interface in $U1$ use only operation1 and $U2$ use only operation2 and $U3$ use only operation3



In a statically typed language (e.g., Java), any change to $op1()$ will force recompilation and redeployment of $User2$ and $User3$.

Is different from the SRP because we need the classes but we need to modularize better.


Inserting interfaces the classes use the interfaces instead the class and the interface implement the class.

If we change operation1 we impact also user2 and user3, if we change the interface we impact only the corresponded user.

We also package the interface and the user together and this not impact the module.

Operations are segregated into interfaces.

In a statically typed language (e.g., Java), a change to `op1()` will NOT force recompilation and redeployment of User2 and User3.



Remove useless operations from modules, otherwise they will force useless redeployments

4.2.3 Dependency inversion principle - DIP


You should always have dependencies that lower modules depend on higher modules, not the opposite.

High-level modules should not depend on anything from low-level modules.

Stable abstractions: abstractions should not depend on details, so that changes to details will not affect them.

Idealistic tension, in practice each system has at least a concrete component: the main function

We can make an interface:



Example In Java (or other statically typed languages), this means that use, import and include statements should refer only to source modules containing interfaces, abstract classes or any other type of abstraction.

In dynamically typed languages, source code dependencies should not refer to modules in which functions being called are implemented.


Also we can package the app and the interface together.

4.2.4 DIP: practical implications

Dependencies on concretions are tolerable when it comes to operating systems, platform facilities and other similar components whose changes are tightly controlled and not frequent.

Abstract Factories are helpful to adhere to the DIP by minimizing the number of dependencies on concrete implementations and keep them separate from the rest of the system


In general, volatility should be shifted towards implementation details without changing interfaces.



- Don't refer to volatile concrete classes
- Don't refer from volatile concrete classes
- Don't override concreted functions
- Never mentions the name of anything concrete and volatile

Using Factory Pattern

To create an instance of `ConcreteImpl`, the `Application` calls `makeService()` from `IServiceFactory` interface, which is implemented by `ServiceFactoryImpl` class. `ConcreteImpl` is created and returned as a `IService`.



Separation abstract and concrete levels: All source code dependencies point toward the abstract level. The flow of control crosses the separation line in the opposite direction: this is the sense of "dependency inversion".

DIP violation DIP violations cannot be entirely removed. Some dependencies on concrete implementations are necessary. They must be minimized and kept separate

How to compile and see the UML `mvn -f .\Path\to\pom_file\pom.xml clean compile site`

Then Java Project > Target > site > index.html

Chapter 5

Components


The logical components are the building blocks of the system. They usually manifest through a namespace or a directory structure, where leaf nodes containing source code represent the logical components.

Logical architecture: the system's logical components and how they interact each other.

Physical architecture: includes also physical artifacts, defining architectural

5.1 Identify components

Example: SPG = Solidarity Purchasing Group, Solidarity towards producers valuing (small vs large, local and seasonal vs global and specific vs mass). Solidarity among clients: Not a normal shop but a group of people buying together, sharing the principles to select producers. This is a software application for **Clients**, **Farmers** and **Shop employees**.



The identification of components is fundamental activity, necessary either for drafting architecture from scratch or reverse-engineer it. It is an iterative refinement process and three approaches can be used for initial identification:

- **Best guess:** based on the system core functionalities
- **Workflow:** Based on the "happy path"
- **Actor/action:** Based on major actions a user can perform

SPG requirements A neighborhood association manages a Solidarity Purchasing Group (SPG). Every Monday, a farm (rotating among a pool of reference suppliers) submits to the SPG a list of available products, their quantities, and their prices.

Once this list is provided, the data is reviewed and confirmed by the SPG employee.

SPG members receive an email with the list of available products and can place their orders either via the website or by visiting the association's headquarters.

When ordering, SPG members specify the types of vegetables they want and the quantities (expressed in kilos, packs, or boxes depending on the type of fruit or vegetable ordered). If the order is placed at the association's headquarters, the member pays immediately.

At the end of the week, the farm delivers the products to the association's headquarters. Shop employees record the delivery, and members receive a notification when their ordered goods arrive. At this point, members can go to the association to collect the package (paying at the same time, if not done at the time of ordering).


Workflow approach

Macro functionality	Component
Insert products	Product submission
View products	Product browser
Order products	Order placement
Buy products	Order Payment

Step	Component
1. Farmer submits available products, employee reviews	Product submission
2. SPG members/clients receive list of products	Client notification
3. Client browse the catalog of product	Product browser
4. Client places an order	Order placement
4a. Employee places an order for the client	Order placement
5. Client pays the order online	Order Payment
5a. Employee submit order payment	Order Payment
6. Employees register delivery of goods	Product delivery
7. Clients are updated on goods' delivery	Client notification
8. Clients retrieve goods	Product delivery

Actor/action

Actor	Action	Component
Client	Search for products	Product browser
	View details about products	Product browser
	Place order	Order placement
	Pay order	Order payment
Farmer	Insert products availabilities	Product submission
Shop employee	Review products availabilities	Product review
	Insert client order	Delegate order
	Insert client payment	Delegate payment
	Register product delivery	Order delivery
System !	Send email notifications	Email Notification

**5.1.1 Putting components together**

- **Tactic level - design:** Components principles: cohesion, coupling
- **Strategic level - architecture:** Component topology, technical partitioning vs domain partitioning, and Physical architecture, Monolithic vs Distributed architecture.
- **And more:** Deployment, communication style and data topology

5.2 Component principles**5.2.1 Component cohesion**

How to put classes inside a component?

Cohesion is the degree to which the classes within a component belong together.

High cohesion means that the classes within a component are closely related in terms of functionality and purpose, while **low cohesion** indicates that the classes are more unrelated and serve different purposes.

It is suggested to follow principle:

- The **Reuse/Release Equivalence Principle - REP**: The granule of reuse is the granule of release.
- The **Common Closure Principle - CCP**: Classes that change together are packaged together.
- The **Common Reuse Principle - CRP**: Classes that are used together are packaged together.

Reuse/Release Equivalence Principle - REP

To support reuse, classes and modules that make up a component must: belong to a coherent group and be released together.

The release process shall produce appropriate notifications and release documentation, so that users can decide whether to integrate new releases.

Common Closure Principle - CCP

Group into components classes that change together, i.e. for the same reason and at the same time.


Separate into different components classes that have different changes processes and motivation.

Note: There is a relationship between the Common Closure Principle (CCP). It is similar to SRP and it addresses the "closure" aspect of OCP

Common Reuse Principle - CRP

It help deciding which classes/module to place into the same components. In such components, usually classes have many dependencies on each other
Complementary perspective:

- users of a component should not depend on unnecessary features;
- if a class isn't closely connected to another class, it shouldn't be in the same component.



Similar to interface segregation design principle (ISP).

Note: Similar to interface segregation design principle (ISP)

5.2.2 Component coupling

- The **Acyclic Dependencies Principle** (ADP): The dependency graph of packages must have no cycles.
- The **Stable Dependencies Principle** (SDP): Depend in the direction of stability.
- The **Stable Abstractions Principle** (SAP): Abstractness increases with stability.


Acyclic Dependencies Principle - ADP

Historical roots on typical issues of team-based development: the *morning after syndrome*¹ and the *weekly build*.

Solution to the problem: partition the development environment into releasable components

Each component has its own responsible developer/team so the releases are uniquely identified and components follow their own, private, development path. E.g. in git versioning system, using separated branches or even separated repositories.

- **Advantages:** Changes made to one component do not affect other teams, until they decide to adapt to the new component's release. Integration happens in small steps
- **Implications:** teams must monitor, manage and, maintain the dependency structure over time. there can be no "cycles" or circular dependencies in the dependency graph.



Example: directed acyclic graph Starting at any component, it is not possible to get back to the origin by following the edges (dependencies).

What happens when a change is made to Presenters? View and Main will be affected and tests with Interactor and Entities.

The bottom-up "building" for system release: Entities -> Database and Interactor -> Presenters, View, Controllers, Authorizer -> Main

Example 2: dependency cycle Now imagine a new dependency arises: User class in Entities uses the Permissions class in Authorizer.

Testing Entities requires now also Authorizer and Interactor.

There are different ways to solve this problem:

- Solution A: Use dependency inversion principle (DIP)
- Solution B: Create a new component on which both Entities and Authorizer depend. Move the class(es) they both depend on into this new component.

¹The morning after syndrome occurs when code worked on, tested, and left working one day is broken the next morning due to changes made by others

5.2.3 The Stable Dependencies Principle - SDP

Depend in the direction of stability. Stability is not about the frequency of changes, but about the amount of work required to make a change.

A cupboard is stable because it takes a lot of effort to move it, while a cup is not stable because it takes little effort to move it.

A component with lots of incoming dependencies is very stable because it requires a great deal of work to reconcile any changes with all the dependent components.

Modules that are designed to be harder to modify should not depend on modules that require little effort to change.

A simple measure of (in)stability

X is very stable because three components depend on it and it depends on no other component, while Y is very unstable because no component depends on it and it depends on three other components.


Instability: goes from 0 to 1, where 0 is the most stable and 1 is the most unstable.

$$I = \frac{\text{Fan-out}}{\text{Fan-in} + \text{Fan-out}}$$

Where

- Fan-in is the number of incoming dependencies
- Fan-out is the number of outgoing dependencies

Example: $I(C_B) = 0.33$. Fan-in = 4 and Fan-out = 2. $I(C_D) = \frac{2}{4+2} = 0.33$.




SDP and stability measure

Depend in the direction of stability.

The $I(\text{component})$ should be larger than the $I(\text{the components that it depends on})$. I metric should decrease in the direction of dependency. $I(C_B) = I(C_C) > I(C_A) > I(C_D)$

SDP violation as system evolves: example Requirements: C should be stable. D should be flexible.



5.2.4 Stable Abstractions Principle - SAP

A component should be as abstract as it is stable. Code that represent high-level policies of the system should be placed into stable components.

SAP, OCP, SDP and abstract classes

Open Closed principle suggests creating classes that are flexible enough to be extended without requiring modification. Abstract classes are best suited to this need. Thus, applying SDP: stability goes into the direction of abstraction


A simple measure of abstractness this ratio:

$$A = \frac{\text{Number of abstract classes and interfaces in the component}}{\text{Total number of classes in the component}}$$

where:



- $A = 0$ means that the component is completely concrete. No abstract classes.
- $A = 1$ means that the component is completely abstract. Only abstract classes.

Note: Distance from the main sequence: $D = |A + I - 1|$



5.3 Strategic decisions

Iterative Refinement



Operational Architectural Characteristics: examples

- **Scalability:** the system's ability to perform and operate as the number of users or requests increases
- **Elasticity:** the system's ability to perform and operate with peaks of requests
- **Responsiveness:** the system's ability to complete assigned tasks within a given time
- **Reliability/safety:** Whether the system needs to be fail-safe, or if it is mission critical in a way that affects lives. Fault tolerance does the software operate as intended despite hardware or software faults.
- **Availability:** How much of the time the system will need to be available; if that's 24/7, steps need to be in place to allow the system to be up and running quickly in case of any failure.
- **Continuity:** The system's disaster recovery capability.
- **Robustness:** The system's ability to handle error and boundary conditions while running, for example, if the internet connection or power fails.

Engineering Architectural Characteristics

- **Maintainability:** the degree of effectiveness and efficiency to which developers can modify the software to improve it, correct it, or adapt it to changes in environment and/or requirements.
- **Testability:** how easily developers and others can test the software? How complete are tests?
- **Deployability:** ease, frequency and overall risk of deployment
- **Evolvability:** system's ability to survive changes in its environment, requirements and implementation technologies

5.3.1 Conway's law

Organizations which design systems...are constrained to produce designs which are copies of the communication structures of these organizations.

The structure of the design will reflect how people communicate and work together




Figure 5.1: Technical partitioning vs domain partitioning

5.3.2 Monolithic versus Distributed architecture

Monolithic: single deployment unit of all code

Distributed: multiple deployment units connected through remote access protocols

Monolithic	Distributed
Layered architecture	Service-based
Pipeline architecture	Event-driven
Modular monolith	Space-based
Microkernel	Service-oriented
	Microservices

Chapter 6

Architecture visualization: C4 notation

6.1 Issues in sw architecture visualization

There are not only one diagrams to visualize software architecture, but many of them. Variety of goals, stakeholders, focus, etc.

- Purpose of elements (boxes, arrows, etc.)
- Relationship between elements
- Levels of abstraction
- No context or logical starting point
- Consistency between diagrams
- Require explanations or descriptions to be fully understood

Diagrams are like maps: different providers, google or openstreetmap have different purposes. For this reason, we will introduce the C4 model, which is a standard for software architecture visualization.

6.2 C4 Model

With this model we can zoom-in and zoom-out to navigate through the different levels of abstraction to use with different stakeholders.

A **software system** is made up of one or more **containers**, each of which contains one or more **components**, which in turn are implemented by one or more code elements. And **people** use the software systems that we build.

It is organized in 4 levels of abstraction:

- Software System
- Container
- Component
- Code Element




Figure 6.1: Static structure

6.2.1 Container

A boundary inside which some code is executed or some data is stored. Containers are **separately deployable**.

Examples: Application (server-side or client side), Database, file system, etc., Microservice (if not a system per se), Mobile app, Script (python, shell, R, etc.)

6.2.2 Component

A grouping of related functionality encapsulated behind a well-defined interface. Components are **NOT** separately deployable and how components are packaged is not relevant at this level.


6.2.3 Code

The basic building blocks of a programming language: classes, interfaces, functions, objects, enums

6.3 Different type of diagrams

The C4 model defines different types of diagrams for each level of abstraction.

- **Context diagram:** it shows how the software system in scope fits into the context around it.
- **Container diagram:** it shows the high-level technical building blocks (containers) and how they interact.
- **Component diagram:** it shows the components inside each container
- **Code diagram:** it shows how a component is implemented



Using the SPG, Solidarity Purchasing Group, presented before, as an example, we can create different diagrams for each level of abstraction.

6.4 Level 1: Context diagram

It helps answering such questions: What is the software system that we are building (or have built)?, Who is using it?, How does it fit in with the existing environment?

It is **required**: it is the starting point for communicating with any stakeholder. The key elements of a context diagram are: **People** who will use the software system, Other **software** systems with which the reference system interacts, + optionally the **organization boundary**.

Interaction between elements: Directed arrows with high-level information about the purpose of that interaction. We can connect elements with arrows. It is better to use only one way arrows and not to use bidirectional ones.

6.4.1 Information reported - People


The people element have:

- **Name:** the name of the person, user, role, actor or persona.
- **Description:** A short description of the person, their role, responsibilities, etc.

6.4.2 Information reported - Software system

The software system element have:

- **Name:** The name of the software system.
- **Description:** A short description of the software system, its goal and responsibilities, etc.



6.5 Level 2: Container diagram


Goal: illustrate the high-level technology elements and how responsibilities are distributed across elements and how containers communicate with each other

Every container must be deployable independently and the diagram is **required** and read by software development and operations staff.

6.5.1 Elements

People, software systems and containers. For each container:

- Name: The name of the container (e.g. “Internet-facing web server”, “Database”, etc).
- Technology: The implementation technology (e.g. Java/Spring MVC application, ASP.NET web application, Relational Database Schema, etc.) or the general type of technology”
- Description: A short descriptive statement, or perhaps a list of the container’s key responsibilities or entities/tables/files/etc. that are being stored.





System boundary must be coherent with L1

6.5.2 Interactions between containers

Interactions between containers must be as explicit as possible.

Example of useful information to annotate the interactions with:

- The purpose of the interaction (e.g. “reads/writes data from”, “sends reports to”, etc).
- The communication mechanism (e.g. REST, Web API, Java Message Service, Web Services).
- The communication style (e.g. synchronous, asynchronous, batched, twophase commit, etc).
- Protocols and port numbers (e.g. HTTP, HTTPS, SOAP/HTTP, SMTP, FTP, RMI/IIOP, etc)



6.6 Level 3: Components diagram

It shows components that are within a single container at a time. Infrastructural elements (e.g. , logging, metering) and shared library might be annotated, separated, or excluded.

It helps to answer the questions like

- What components is each container made up of?
- Do all components stay within in a container?
- How components interact each other ?
- How the software works at a high-level?

Optional, especially for containers like data storage, microservices. It is for technical people within the software development team.

Chapter 7

Writing Components

7.1 Components vs. Service

- **Component:** reusable piece of software that is used locally, eg static/dynamic linking. Is out of control of the writer of the component.
- **Service:** reusable piece of software that is typically accessed remotely, either synchronously or asynchronously, eg web service, RPC or socket.

7.2 Framework

A framework is a library that provides basic reusable assets that can be extended by the developers. Instead of writing from scratch, use some code that already exists.

It provides an abstraction layer over lower-level code, allowing developers to focus on higher-level features.

Note: Extension and customization are often based on the inversion of control.

7.2.1 Inversion of control

"Don't call us, we'll call you"

With inversion of control, your code does not control the overall flow of the program. Instead, a framework or container controls execution and calls your code at the appropriate time.

Example of direct control

```
void main() throws IOException{
    Path path = Path.of("data/grades.csv");
    List<String> lines = Files.readAllLines(path);
    List<String> h = Arrays.asList(lines.get(0).split(","));
    processHeaders(h);
    lines.stream().skip(1)
        .map(r -> Arrays.asList(r.split(",")))
        .forEach(r -> processRow(r));

    double average = sum / (lines.size()-1);
    IO.println(average);
}
```




Figure 7.1: Call diagram

Example of inverted control

```
void main() throws IOException{
    CsvReader reader = new CsvReader();
    reader.headerProc(ReadCsvIoC: :processHeaders);

    reader.rowProc(ReadCsvIoC: :processRow);

    int n = reader.read("data/grades.csv");
    IO.println(sum/n);
}
```




Figure 7.2: Call diagram

7.2.2 Patterns that enable inversion of control

Template method - Strategy pattern - Observer pattern

7.3 Wiring components

7.3.1 Interaction mode

- **Call-based:** function or method calls, the base logic of lots programming languages.
- **Event-based:** messages are sent and received. Used in GUI and web applications.
- **state-based:** state is updated and read from a shared memory. Used in embedded systems, writing in a specific memory address, and some peripheral devices read it.

7.3.2 Call binding

- **Source inclusion:** The code of the component is included directly in the main program during compilation. Strong implementation and compile-time dependencies.
- **Static linking:** Strong implementation dependency, limited compile time.
- **Dynamic linking:** Limited implementation dependency, limited compile time
- **Dynamic loading:** Strong construction dependency, no compile time


7.3.3 Method/function calls

Problem of simple calls:

- Structural dependencies
- Runtime dependencies: spacial, temporal and synchronization dependencies
- Cascaing failures
- Limited scalability
- Hard to evolve independently


How put toghere components while minimizing the dependencies? E.g. they are build by different teams with little knowledge of each other.

Example A *ThesisLister*, frontend callses, provides a list of thesis matching a keykords using *ThesisFinder*.



Is not testable, not replaceble, heavy dependencies.

Using **interface** instead of concrete implementation to decouple the components, but still have dependencies.



7.4 Decoupling solutions

To solve the problem of dependencies, we can use 3 techniques for decoupling components :

Dependency injection, Service locator, Message

7.4.1 Dependency injection


The *ThesisLister* calasss id dependent on both the *ThesisFinder* interace and the implementation of *ThesisFinderImpl*.

The goal of dependency injection is make the lister class unaware of the implementation class, but can still talk to an instance to do its work.

Move the logic of creating into an *Assembler* class

How to performe the injection? Two techniques:


- **Constructor injection:** use the constructor of the dependent class to inject the dependencies.
- **Setter injection:** objects are initialized with dependency using a setter



7.4.2 Constructor injection

Basic injection

```
ThesisFinder finder = new ThesisFinderImpl(`theses.csv`); //creates the
↳ implementation with parameters
ThesisLister lister = new ThesisLister(finder); //injects the dependency
↳ upon construction
```



Constructor injection with container

```
Container c = new Container();
Object[] parameters = {"theses.csv"}; //parameters for the implementation ctor
c.registerImplementation(ThesisFinder.class, ThesisFinderImpl.class, parameters); // binding interface to
↳ implementation
c.registerComponent(ThesisLister.class); // component registration
```

Reflection is used to find out how to link.

Constructor injection with annotations

```
@Component
public class ThesisFinderImpl implements ThesisFinder {
    public ThesisFinderImpl(@Value("${thesis.file.path}") String file) {
        this.file = file;
    }
}

@Component
public class ThesisLister {
    public ThesisLister(ThesisFinder thesisFinder) {
        this.thesisFinder = thesisFinder;
    }
}
//...
```


Note: The Spring framework is based on constructor injection with annotations.

7.4.3 Setter injection

Instead of using the constructor we can use a setter method to inject the dependencies.

Basic injection

```
ThesisFinderImpl finder = new ThesisFinderImpl(); //creates the server
↳ object
finder.setFile("theses.csv"); //set the parameter
ThesisLister lister = new ThesisLister() //creates the client
lister.setFinder(finder); //injects the dependency with the setter
```



To set up a setting injection we can use a configuration file.

Setter injection with configuration file

```
<beans>
  <bean id="ThesisLister" class="it.polito.
    ↳ ThesisLister">
    <property name="finder">
      <ref local="ThesisFinder"/>
    </property>
  </bean>
  <bean id="ThesisFinder" class="it.polito.
    ↳ ThesisFinder">
    <property name="file">
      <value>theses.csv</value>
    </property>
  </bean>
</beans>
```

```
[
  {
    "id": "ThesisLister",
    "class": "it.polito. ThesisLister",
    "properties": [
      { "name": "finder", "local": "ThesisFinder" }
    ]
  },
  {
    "id": "ThesisFinder",
    "class": "it.polito. ThesisFinderImpl",
    "properties": [
      { "name": "file", "value": "theses.csv" }
    ]
  }
]
```

Setter injection with annotations

```
@Component
public class ThesisLisiter {

    private ThesisFinder finder;

    public ThesisLisiter() {}

    @Autowired
    public void setThesisFinder(ThesisFinder thesisFinder) {
        this.thesisFinder = thesisFinder;
    }
}
```

Criterion	Constructor	Setter
Mandatory	Yes (object creation)	No (later stage)
Validity	Yes, always	Only after injection
Immutability	Yes (final)	No
Null safety	Fails fast	Risk later <code>NullPointerException</code>
Understandability	Depends explicit ctor	Depends hidden in setters
Circular dependencies	No	Can sometimes resolve
Optional dependencies	No, requires <code>Optional<></code> or ctor overload	Yes (with <code>@Autowired(required = false)</code>)

7.5 Service locator


A service locator is an object that knows how to get hold of all of the services that an application might need.

It is an alternative to dependency injection to break the dependencies.

Example of locator service

```
public class ThesisLisiter {
    public static String FINDER_SVC="thesis-finder";
    private final ThesisFinder thesisFinder;

    public ThesisLisiter() {
        this.thesisFinder = Locator.getLocator().lookup(FINDER_SERVICE,
            ThesisFinder.class);
    }
    //...
}
```



Basic assembler


```
Locator l = Locator.getLocator();
l.register(ThesisLisiter.FINDER_SVC, new ThesisFinderImpl("theses.csv"));
ThesisLisiter thesisLisiter = new ThesisLisiter();
```

7.5.1 Locator vs. Dependency injection

Criterion	Service Locator	Dependency Injection
Decoupling	Yes	Yes
Mechanism	Client explicit lookup	Provided externally
Inversion of Control	No	Yes
Visibility of dependencies	Hidden inside method calls to locator	Explicit in constructor/setters
Dependency on infrastructure	Client depends on the locator	No dependency on injector/container API
Ease of debugging	Easier (explicit calls)	Harder (framework wiring)
Simplicity	Straightforward	Less obvious at first
Reusability	Limited	Yes, no assumptions
Deferred linking	Yes, call locator anytime	No, injected at startup

7.6 Messages/Events

Replace Direct call with a Broker



In this case sender does NOT know: Who consumes the message, When the message is processed, How many consumers exist

7.6.1 Roles: Producer, Consumer

Producer	Consumer
Sends messages	Receives messages
Does not wait for processing	Processes independently
Does not know consumers	Can scale horizontally

7.6.2 Roles: Broker

The broker, like in real life, is an intermediary that facilitates the communication between producers and consumers. It's a middleware responsible for:

- Storing messages
- Delivering messages to consumers
- Handling routing and filtering of messages

Some examples of brokers are RabbitMQ, ActiveMQ, Kafka, Amazon SQS.

7.6.3 Delivery Guarantees

It is possible to configure different levels of delivery guarantees depending on the requirements of the application.

Type	Meaning
At most once	messages may be lost
At least once	messages may be duplicated, need idempotency
Exactly once	delivery guarantees, expensive

Distributed systems always involve trade-offs.

Producer Example

```

ConnectionFactory factory = new ConnectionFactory();
factory.setHost("localhost");

try (Connection connection = factory.newConnection();
    Channel channel = connection.createChannel() ) {
    channel.queueDeclare("orders", false, false, false, null);
    String message = "New Order: 123";
    channel.basicPublish("", "orders", null, message.getBytes());
    System.out.println("Sent: " + message);
}

```

Producer Example

```

ConnectionFactory factory = new ConnectionFactory();
factory.setHost("localhost");
Connection connection = factory.newConnection();
Channel channel = connection.createChannel();
channel.queueDeclare("orders", false, false, false, null);
DeliverCallback callback = (tag, delivery) -> {
    String message = new String(delivery.getBody());
    System.out.println("Received: " + message);
};
channel.basicConsume("orders", true, callback, tag -> {});

```

7.6.4 Structure

Producer and consumer are independent processes, they can run on different machines and the broker is infrastructure.

7.6.5 Shared state example

Queue	Topic
1 message → 1 consumer	1 message → many consumers
Work distribution	Event broadcast
Load balancing	Publish/subscribe

7.6.6 Effects

- **Construction decoupling:** producer does not create consumer
- **Spatial decoupling:** producer does not know location of consumer
- **Temporal decoupling:** consumer does not have to be active with producer. Producer does not have to block waiting for an answer
- **Data dependency:** data format and schema must be shared

7.6.7 Challenges

- Difficult debugging
- Eventual consistency
- Duplicate handling (idempotency!)
- Ordering issues
- Schema evolution

7.7 Shared state

Is a database just storage, or is it a communication mechanism?

Communication via mutations of shared, persistent state. The components communicate by reading and writing a common data store, no explicit message passing.

Shared state examples

- In-memory global variables
- Shared files
- Shared database
- Distributed cache (Redis)

Communication patterns

- State-based signaling
- Polling
- Work queue
- Derived state sharing

7.7.1 Effects

- No structural dependencies
- Low temporal dependencies: DB is like message broker
- Very high data dependency
- Simple mental model
- Easy querying and debugging

Aspect	Shared DB	In-memory Cache (Redis)
Nature	Persistent system of record	Ephemeral shared memory
Latency	Higher	Very low
Consistency	Strong	Config-dependent
Visibility	High (queryable history)	Low (state disappears)

Chapter 8

Communication

The communication between components can be implemented in different ways:

- Local vs Distributed
- Synchronous vs Asynchronous
- One-to-one vs One-to-many
- One-way vs Two-way
- Call vs Message vs Shared memory

There are some distributed computing fallacies that are worth mentioning:

1. The network is reliable;
2. Latency is zero;
3. Bandwidth is infinite;
4. The network is secure;
5. Topology doesn't change;
6. There is one administrator;
7. Transport cost is zero;
8. The network is homogeneous.

Network latency - Due to decomposition an operation may result into a series of round-trips between two services. Some solution can be: Batch API that allow feeding multiple requests or merge services to convert messages to method calls,

	Multiplicity	
	One-to-one	One-to-many
Sync	Request/response	–
Async One-way	Notification	Publish/Subscribe
Async Two-way	Async. Req./Resp.	Publish/Async. Resp.

8.0.1 One-to-one communication

- **Request-response** - A client makes a request to a service and waits for a response. The client blocks while waiting.
- **Asynchronous request/response** - A client sends a request to a service, which replies asynchronously. The client doesn't block while waiting.
- **One-way notifications** - A client sends a request to a service, but no reply is expected or sent.

8.0.2 One-to-many communication

- **Publish/subscribe** - A client publishes a notification message, which is consumed by zero or more interested services.
- **Publish/async. response** - A client publishes a request message and then waits for a certain amount of time for responses from interested services.

8.1 Messaging

Messages are exchanged over message channels. A sender writes a message to a channel, and a receiver reads messages from a channel. The communication is asynchronous, the sender doesn't block waiting for a reply

Header The header contains metadata about the message, such as: Message ID, Return address, Metadata that describes the content of the message, Collection of key-value pairs useful for routing, filtering, etc.

Body The body contains the data being sent in either text or binary format.

8.1.1 Message types

- **Command** - This is equivalent of an RPC request. It specifies the operation to invoke and its parameters.
- **Document** - Contains only data. The receiver decides how to interpret it, e.g. a reply to a command
- **Event** - Indicates that something notable has occurred. An event is often a domain event, which represents a state change of a domain object.

8.2 Channel


8.2.1 Point-to-point channels

Delivers a message to exactly one of the consumers that is reading from the channel.

Services use point-to-point channels for the one- to-one interaction styles. E.g. command messages

8.2.2 Publish-subscribe channels

Channel delivers each message to all of the attached consumers. Services use publish-subscribe channels for the one-to-many interaction styles



8.3 Supported styles

- Request/response - supports synchronous or asynchronous
- One-way notification - No response
- Publish/subscribe - Multiple subscribers no response
- Publish/Async responses - Multiple subscribers, multiple responses

8.4 Direct messaging - brokerless

Pros	Cons
Lighter traffic load, less latency	Services must be known to discovery
No centralized point of failure	Reduced availability, both must be available at the same time
Simplicity	Hard to have guaranteed delivery

8.5 Broker-based messaging

Broker is an intermediary components through which all messages flow.

Pros	Cons
Loose coupling.	Potential performance bottleneck
No need to know each other or to be active at the same time	Single point of failure
	Additional complexity

8.5.1 Broker characteristics

It support both standard protocols, e.g. AMQP, MQTT, STOMP, or proprietary protocols. It has some good feature such as:


- Messaging ordering
- Guaranteed delivery
- Persistence
- Durability, consumer can disconnect and reconnect without losing messages
- Scalability
- End-to-end latency
- Support for cometing consumers

8.6 Coordination

8.6.1 Choreography

Each local transaction publishes domain events that trigger local transactions in other services.

- It requires reliable message infrastructure and messages with correlation id.
- Basic operation is simple, it perform operation and publish event
- Difficult to understand and maintain because transaction logic is scattered
- Complex service dependencies also cyclic, might lead to tight coupling





8.6.2 Orchestration

An orchestrator (object) tells the participants what local transactions to execute.

This technique requires reliable messaging, it has a simple dependency structure with reduced coupling.

The business logic is simple and services are not aware of the saga. Transaction logic is gathered in a single place, must be kept separate from service logic.




Chapter 9

Microservices

The scalability dimension of software architecture can be represented by the following figure:


Advantages	Disadvantages
Enables the continuous delivery and deployment of large, complex applications	Finding the right set of services is challenging
Services are small and easily maintained	Distributed systems are complex, which makes development, testing, and deployment difficult
Services are independently deployable	Deploying features that span multiple services requires careful coordination
Services are independently scalable	Deciding when to adopt the microservice architecture is difficult
Enables teams to be autonomous	
Allows easy experimenting and adoption of new technologies	
Better fault isolation.	



9.0.1 Monolith vs. Microservices

The problem is the deployment in fact:

- Monolithic architecture: Simple to develop, deploy, and scale.
- Microservices architecture: Simple to maintain, to test and flexible to scale




9.0.2 Decomposition


Decomposition identifies system operations:

- Understand domain model
- Define operations
- Identify use cases

and decompose the system into subsystems or subdomains.



Domain model Example



User stories


Actor	Goal	Services
Consumer	Create order	Order
Restaurant	Accept order	Order
Restaurant	Complete order	Kitchen
Courier	Update location	Delivery
Courier	Pick up	Delivery
Courier	Delivery	Delivery
Delivery company	Manage couriers	Courier
Delivery company	Courier payment	Accounting

By business capability

1. An organization business capability is what it does (as opposed to how). Look into organization structure, also consistent operations, stories or goals.
2. Single Responsibility Principle
3. Common Closure Principle - classes that change for the same reason should be in the same package
4. Result is cohesive and loosely coupled

By subdomain

- Identify bounded contexts - DDD sub-domains. Partitions of overall domain model
- Result is cohesive, Loosely coupled, allow cross-functional and autonomous teams and scalable architecture



9.0.3 Guidelines from OOP

Single Responsibility Principle - A class should have only one reason to change.

Common Closure Principle - The classes in a package should be closed together against the same kinds of changes. A change that affects a package affects all the classes in that package.

Decomposition obstacles


: Network latency, Reduced availability due to synchronous communication, Data consistency across services, Obtaining a consistent view of the data, Good classes

9.1 Business logic

9.1.1 Transaction script

When you have simple business logic that consist of a sequence of steps. Organize business logic by procedures where each procedure handles a single request from the presentation. One service class with a method for each procedure. Several data classes. Consequences:

- All the code for a request in a single place
- Code duplication for similar requests
- Difficult to maintain and refactor



9.1.2 Domain model

When you have complex business logic

Adopt recommended object-oriented methods to create classes that encapsulate both code and data, using the DDD approach.

Consequences:

- Code for a request is spread among classes
- No code duplication
- Easy to maintain and refactor
- Easy to extend with Strategy or Template Method patterns




Figure 9.1: Ports & Adapters (hex)

9.2 Domain Driven Design - DDD

- Sub-domains
- Bounded context
- Class types patterns: Entity, Value Object, Service, Repository, Factory
- Aggregate

9.2.1 Sub-Domain & Bounded Context

DDD decompose the overall domain model into sub-domains. Each sub-domain is valid within a bounded-context.

A bounded-context brings together a common language and people, team, focusing on a shared problem. Distinct contexts are related to each other in order to enable communication.


9.2.2 Class patterns

- Entity - Represents objects that have a persistent identity. Two objects with the same values represent distinct items. Usually, they are persisted to a db
- Value object - Represents a collection of values. Two objects with the same values are interchangeable
- Factory - Implement object creation logic. May hide concrete class
- Repository - Provides access to persistent entities. Hides access to db
- Service - Implements business logic that does not belong to a single entity

9.2.3 Domain boundaries

Often domain models have fuzzy boundaries. All strictly related concepts should belong to the same context. E.g. **Order** and **OrderLineItem** classes.

Not keeping together such classes might result in violations of business rules or invariants, upon updates. E.g. minimum amount for an order.




9.2.4 Aggregate pattern

Organize a domain model as a collection of aggregates, each of which is a graph of objects that can be treated as a unit.

Decompose domain model into chunks:

- Self contained units
- Easier to understand
- Clarify scope of operations
- All elements share the same life cycle



9.2.5 Rule: reference roots only

Reference only the aggregate root. Updates go through methods of aggregate root. When a repository loads an aggregate from the db it returns a reference to the aggregate root.

Consequences:


- Enable aggregate to enforce consistency rules
- Aggregate updates are serialized

9.2.6 Rule: primary keys inter aggregate

Inter-aggregate references **must use primary keys**. NOT using object reference, indeed, this is usually code smell.

Consequences:

- Loose coupling of aggregates
- Keep aggregate boundaries enforced
- Avoid accidental cross-aggregate updates
- Aggregates can belong to distinct services
- Simplify persistence




9.2.7 Rule: one aggregate per transaction

One transaction operates on one aggregate only


Consequences:

- Ensure a transaction is contained in a single service
- Suitable for limited transactions support in NoSQL databases
- Less efficient in transfer from/to db
- More complex implementation of complex requests that update multiple aggregates




9.3 Split monolith with aggregates

Transformation into aggregates can help splitting the monolith. Services can be extracted easily.



9.3.1 Aggregate granularity

- Fine grained: Improve scalability (Updates on an aggregate are serialized), reduce the risk of conflicts and it is easier to decompose into services
- Coarse grained: More suitable to fit all relevant entities in a transaction



9.3.2 Domain Events

An Aggregate publishes a Domain Event when it's created or undergoes some other significant change. Represented by a class in the domain model, the name is formed using a past-participle verb phrase. E.g. **OrderPlaced**, **OrderShipped**, etc.


Properties that convey the event meaning. Each property is either a primitive value or a value object, also metadata, e.g., event ID, timestamp

- Maintaining data consistency across services
- Notifying a service that maintains a replica that the source data has changed
- Sending notifications to users
- Notifying another service to update related data
- Notifying a different application via a registered webhook or via a message broker
- Monitoring domain events to verify that the application is behaving correctly.
- Analyzing events to model user behavior

9.4 Communication

Ensure agreement on service interaction mechanism to avoid late integration issues:

- Define the interface (using an IDL)
- Review with client services' developers
- Iterate to reach consensus
- Start implementation of services



9.4.1 API evolution

"An interactive software system must be continually adapted or it becomes progressively less satisfactory".

Aim for minor backward-compatible changes.

Robustness principle:

- Be conservative in what you do
- Be liberal in what you accept from others

9.5 Message formats


Text-based formats: JSON, XML, YAML

Binary formats: Protocol Buffers, Avro, Thrift

Once the format is chosen, it might be mandated by the communication mechanism

9.6 Remote procedure call (RPC)

- **Pros:** **Simple** and familiar, request-response is easy and the system is simple because of no intermediate broker
- **Cons:** Usually **only supports request/reply** and not other interaction patterns such as notifications, request/async response, publish/subscribe, publish/async response. **Reduced availability** since the client and the service must be available for the duration of the interaction



Requires service discovery

9.6.1 Robust Synchronous APIs

- Network timeouts: Never block indefinitely, always use timeouts when waiting for a response
- Limit outstanding requests: Define an upper bound on the number of outstanding requests that a client can make to a particular service
- Circuit breaker pattern, allow failing fast

9.6.2 Circuit Breaker pattern

Service should be invoked through a proxy. Track the number of successful and failed requests. If the error rate exceeds a threshold, trip the circuit breaker so that further attempts fail immediately. After a timeout period, the client should try again, and, if successful, close the circuit breaker.

This is implemented by Netflix's Hystrix library.


9.6.3 Service Discovery

Using RPC you need to know the network location of your service.

In microservices architectures service instances change continuously due to: Autoscaling, Failures, Updates.

Application level discovery

- Self registration pattern: service register itself on activation and update list (Health check endpoint on server, Heartbeat API callback on registry)
- Client-side discovery pattern: Queries registry for active servers. Uses load balancing algorithm to pick one.



Example: confirm order

- **Order Service:** Create an Order in an APPROVAL_PENDING state.
- **Consumer Service:** Verify that the consumer can place an order.
- **Kitchen Service:** Validate order details and create a Ticket in the CREATE_PENDING state.
- **Accounting Service:** Authorize consumer's credit card.
- **Kitchen Service:** Change the state of the Ticket to AWAITING_ACCEPTANCE state.
- **Order Service:** Change the state of the Order to APPROVED state.

9.9.3 Distributed data consistency


- **Distributed transaction - 2PC:** Query to commit, Agree/Abort, Commit/Rollback, Acknowledge.
- **Saga:** Local transactions coordinated with messages.

Distributed transactions limits

Transactions are not universally supported (e.g. NoSQL dbs, Modern message brokers). Transactions are synchronous, which means reduced availability. Isolation is overrated (most dbs adopt weaker isolation levels). Impossible to have both availability and most recent data.

9.9.4 Saga pattern

A saga is a sequence of local transactions. Each local transaction updates the database and publishes a message or event to trigger the next local transaction in the saga. If a local transaction fails then the saga executes a series of compensating transactions (undo the changes that were made by the preceding local transactions).

**Compensating transactions**

Rollback is not directly supported since all service transactions are separate. Explicit invocation of compensating transactions in case of rollback (Used, e.g., when a check in a step fails).

Saga transactions types

- **Compensatable transaction:** Can potentially be rolled back using a compensating transaction.
- **Pivot transaction:** The go/no-go point in a saga. If succeeds the saga can run until completion.
- **Retriable transaction:** Follow the pivot transaction and is guaranteed to succeed.

Saga non-isolation

- **Atomic:** Managed by coordination mechanism and compensating transactions.
- **Consistent:** each service guarantee consistency in its own transaction.
- **Durable:** ensured by local db.

Non isolated problems: Lost updates, Dirty reads, Non-repeatable reads.

Countermeasures for non-isolation

- **Semantic lock:** An application-level lock. A compensatable transaction sets a flag in any record that it creates or updates indicating it isn't fully committed. It's cleared by either a retrieable transaction or a compensating transaction.
- **Commutative updates:** Design update operations to be executable in any order (e.g., `Account.debit()` decrement balance).
- **Pessimistic view:** Reorder the steps of a saga to minimize business risk due to a dirty read. Cancel order (rolled back) and create order when executed concurrently could lead to exceeding customer's credit.
- **Reread value:** Prevent dirty writes by rereading data to verify that it's unchanged before overwriting it. If changed, saga aborts.
- **Version file:** Record the updates to a record so that they can be reordered (turns noncommutative operations into commutative).
- **By value:** Use each request's business risk to dynamically select the concurrency mechanism (sagas vs distributed transactions).

9.9.5 Domain event pattern

An aggregate publishes a domain event when it's created or undergoes some other significant change. A domain event is a class with properties that meaningfully convey the event, a name formed using a past-participle verb (e.g. `OrderCreated`), and metadata (e.g. event ID, timestamp).

9.9.6 Event sourcing

Event sourcing is an event-centric technique for implementing business logic and persisting aggregates. An aggregate is stored in the database as a series of events, where each event represents a state change. It is an alternative to common persistence, e.g. through ORM.

Pros	Cons
Reliably publishes domain events	It has a different programming model that has a learning curve
Preserves the history of aggregates	Messaging-based application are complex
Avoids the O/R impedance mismatch	Evolving events can be tricky
Provides developers with a time machine	Deleting data is tricky
	Querying the event store is challenging


9.10 Queries

Queries can be Simple (require data from a single μS db) or Composed (required data from multiple μS dbs).

9.10.1 API Composition

- **API Composer:** Implement a query that retrieves data from several services by querying each service via its API and combining the results.
- **Provider Service:** Implement a service-limited query, limited by how data is partitioned and capabilities of APIs/databases.



Composer role can be Client (e.g. web app backend), API Gateway, or a Standalone service. For efficiency, parallelize queries as much as possible.



9.10.2 CQRS

Command - Query Responsibility Segregation. Use events to maintain a read-only view that replicates data from the services. The view can be easily queried.


Design the schema to match the queries. The data access module must handle idempotent updates and concurrent updates. Implement a mechanism to efficiently build or rebuild the view, and cope with the replication lag.




9.11 External API

External clients have low bandwidth (negative impact on UX) and lack of encapsulation. Different communication mechanisms are used. Possible solutions:

Edge Functions Authentication, Authorization, Rate limiting, Caching, Metrics collection, Request logging.



Backend for Frontend (BFF)



9.12 Transition to Microservices

- **Big-bang rewrite:** "the only thing a Big Bang rewrite guarantees is a Big Bang!" (M.Fowler) .
- **Strangling the monolith:** Progressively wrap the monolith until it becomes an empty shell.


9.12.1 Strangling the monolith

Context: a monolith application should be migrated to a micro service one

Problem: a complete big-bang rewrite poses several risks

Solution: gradually create a new system around the edges of the old, letting it grow slowly over several years until the old system is strangled.

Consequences: Reduced risks, frequent releases allow to monitor the process and provide steady release of value, speed of transition can be adjusted.



9.12.2 Transition strategies

- **Implement new features as services:** Stop the growth of the monolith. API gateways reroute requests. Integration glue code integrates the service with the monolith (typically adapters).
- **Separate presentation and backend:** Represents an intermediate strategy.
- **Break up the monolith:** Extract functionalities into services. Extract a vertical slice (Inbound adapters, Domain logic, Outbound adapters, Monolith's database schema).




Figure 9.2: New features as services




Figure 9.3: Split separation and backend

9.12.3 Refactor DB

Several classes are persistent so the relative tables in the db must be refactored. Replicate data to avoid widespread changes in the monolith (Keep original schema and use triggers to update data).


Selecting services to migrate Rank modules by anticipated value your organization can get by migrating:

- Features providing new business capabilities
- Accelerated development
- Performance, scaling, reliability issues
- Unlock migration of other more valuable modules

9.12.4 God Class

An object that controls way too many other objects in the system and has grown beyond all logic to become The Class That Does Everything. Often implements business logic for many different aspects.

Splitting the God Class: Apply DDD to separate domain models into subdomains. Avoid tight coupling among services and avoid a data service. Different views must be kept consistent (via Saga) and UX must be adapted to the set of partial models (via API gateway).



Chapter 10

Bending Spoons seminary (Evernote)

Founded in 2013 in Danmark. They acquire digital products then unlock the potential of the products. 500+ TB data processed per year. Constantly search for new talent.

10.1 Evernote

There is a frontend that get the data from backend at the very high level, but deeper the system is very complex and full of components. Before the acquisition Evernote was a monolith, but after the acquisition they decided to migrate into a set of microservices with each one take care of a specific element. All of these services should minimize the dependencies with the other services. But the backend is not the complex part of the system.

10.1.1 Evernote's sync system

Legacy: "Monolith have you some new content?" "Yes, the note with the monkey" This is a polling request useless for the user and for the backend system.

Microservices: Persistent connection between the client and the server, faster than polling, cheaper for both: the client side save battery and for the server side save CPU and bandwidth.

Deep dive into sync activity

Note creation:

- The phone perform an http request to the backend to note service
- Performing some checks
- Create the json file SQL database
- Note create message to the sync service queue
- The sync service send the message to the all the client devices with the new note

Client connection

Fernando authenticate with the websocket and then connect to sync service. This service has tons of persistent connection always open, it use a multi-map, because of multiple devices but same user, to keep track of the connection and the user id. The persistent connection can go offline for a bit. Do they lose all the message? Yes but the sync service use google cloud datastore. It is a single schema with the UID, the Instance with every metadata of Fernando and SyncOperation with the operation performed (CREATE, UPDATE, DELETE). Every time the client receives some message, it receives the userId and the timestamp and this is used in the google cloud datastore.

```
SELECT * FROM sync_documents
WHERE accessInfo.updated >= "1259414" and accessInfo.userId = "12345"
```

10.1.2 The sync service

Pros	Cons
Fast updates	Data duplication
Offline access	Integration difficulty
Cheap	Complexity

Chapter 11

Boundary

Architecture is defined by the boundaries drawn within a system and the dependencies that cross those boundaries, rather than the physical mechanisms through which components interact and execute functions.. Boundaries between components allow to make “firewalls” against harmful propagation of changes. The components on one side of the boundary change at different rates and for different reasons than the components on the other side of the boundary


11.1 Drawing boundaries

1. Partition the system into components
 2. Identify core business components from those containing functions not related to the core business
 3. Arrange the code so that dependencies point toward the core business
- Make use of SOLID principles and Component principles.

11.1.1 The clan architecture

- Outer circles: mechanisms.
- Inner circles: policies.

Dependency Rule : source code dependencies point only inward, toward inner circles (i.e. higher-level policies)



11.1.2 Independence and Decoupling of layers

Layers allow independent development, deployment, operation, and maintenance. Two type of layers: Horizontal (DB, UI, etc.) and Vertical (Source code units, services, etc.)

Decoupling of policy from the details. Defer decisions on details for as long as possible.

DB component knows about use case rules component, but NOT viceversa. **Low level components know** everything of the **higher-level** components, but NOT viceversa.

The Database component contains the code that translates the calls made by the BusinessRules into the query language of the database. This translation code knows about the BusinessRules. Business rules can be adapted to any kind of DB.




11.2 Crossing boundaries

Are required for: require a service, ask for data or forward requests. Only data cross boundaries, not entities. Data is passed in the most convenient format for the inner circle. Entity objects or database rows are NOT passed, otherwise Dependency Rules might be violated.


Low-level client to high-level server

The Client calls a function $f()$ on the Service. Data is passed as argument of the function. Dependencies cross boundary from left to right point toward the higher-level component.



High-level client to low-level server

The Client calls the $f()$ function of the ServiceImpl, through the Service interface. Dependencies cross the boundary always toward the higher-level component.



11.2.1 Data is significant. Data storage is a detail

From an architectural perspective, the database is NOT an entity. It is a detail that does not rise to the level of an architectural element.

It is an architectural mistake to allow database rows and tables to be passed around the system as objects. This is because it causes the use cases, business rules and, in some cases, also the UI to be coupled to the relational structure of the data.


The the data model, i.e. the organizational structure of data, instead, is architecturally significant.

Example In our university, we need a software that manages the identities. One of the use cases is to check identity from the polito badge and grant/deny access for a specific area or room of the university. It should be possible to check whether a person (student, professor, etc) have access to a given laboratory in each department.

Policy about permissions (including which days and times, and for how long time) are defined on a specific database. Card readers use the magnetic stripe, but soon there will be also NFC card readers in some department accesses. A red or green light on the reader will show the result of the output

Key aspects

What is policy and what details ? Rreading bages, database, rules, output
What is low level and what is high level ?




11.3 Layers

The most basic components of a system are:

- User Interface (UI)
- Application Logic
- Database

Example: a word processor A word processor (WP) is a device or software for creating, editing, storing, and printing text documents.



11.4 Test

Tests support development, not operation, HOWEVER They are an integral component of the system, akin any other elements within the system participate in its architecture. They are part of the design of the system architecture. Also tests adhere to the Dependency Rule.


11.4.1 Design for testability

Tests are coupled to other system's elements. Changes to components in the inner circles may break several tests, to reduce tests fragility, tests should depend as little as possible to volatile elements.

Example: business rules should be tested without using the GUI. Specific tests API can avoid constraints (e.g. security).

11.4.2 Humble objects

The Humble Object helps unit testers to separate behaviors that are hard to test from behaviors that are easy to test. It is adopted in software architecture to identify and protect architectural boundaries.



Example are: Database gateways, Data mappers and Service listeners.

11.5 Architectural styles


Examples of SPG/order.

Use case: customers should be able to view the status of their orders

Classes:

- OrdersController: web controller (e.g. Spring MVC controller) to handle requests from the web.
- OrdersService: interface that defines the logic related to orders.
- OrdersServiceImpl: The implementation of the orders service.
- OrdersRepository: An interface that defines how we get access to persistent order information
- JdbcOrdersRepository: An implementation of the repository interface.

option could be Package by layer or feature or component, Ports and Adapters



11.5.1 Package by layer

Code is organized into layers according to its behavior. Each layer depends only on the next adjacent lower layer. E.g.: webcode, business logic, persistence

- **Pros:** simple, Tied to technical view, Commonly used
- **Cons:** As software grows, layers become “fat”, It is not tied to a the specific “business” view (the same for any domain). Dependencies are still possible between not adjacent layers

11.5.2 Package by feature

Vertical slicing by feature

- **Pros:** It well reflects the domain, Everything about a functionality is in one package
- **Cons:** Suboptimal in terms of separation of implementation details and business logic

11.5.3 Ports and Adapters

Domain-specific code (inner circle) is well separated from implementation details (outer circles). The latter depends on the former: the dependencies flow from outside toward inside.

11.5.4 Package by component

Service-centric view. Everything needed for orders is in the OrdersComponent. Separation of concerns is maintained inside each compone

11.6 Architectural styles

Reusable set of structural principles and high-level design decisions on:

- Component topology: by technical capabilities/by domain
- Physical architecture: monolithic/distributed
- Deployment: system's granularity and deployment frequency
- Communication style: synchronous/asynchronous, protocols/calls/queues
- Data topology: monolithic database/separate data
- Team topologies: remember Conway's law

Monolithic = single deployment unit of all code

Distributed = multiple deployment units connected through remote access protocols

Monolithic	Distributed
Layered architecture	Service-based
Pipeline architecture	Event-driven
Modular monolith	Space-based
Microkernel	Service-oriented
	Microservices

Chapter 12

Monolithic Architectural Styles

12.1 Architectural Quanta

The architectural quanta is the smallest part of the system that runs independently. An architectural quanta establishes the scope for a set of architectural characteristics. It features:

- Independent deployment from other parts of the architecture
- High functional cohesion
- Low external implementation static coupling
- Synchronous communication with other quanta




Figure 12.1: Scoping

12.2 Architectural styles

- Monolithic = single deployment unit of all code
- One architectural quantum
- Distributed = multiple deployment units connected through remote access protocols

Monolithic	Distributed
Layered architecture	Service-based
Pipeline architecture	Event-driven
Modular monolithic style	Space-based
Micro-kernel architecture	Service-oriented
	Microservices


12.3 Layered architecture - n tier

Also known as n-tier architecture. Technical partitioning. Each layer is independent and has a specific responsibility. A layer represents a function of the system. Pro: testability

Common patterns are:

- Presentation layer
- Business logic layer
- Persistence layer
- Database layer

In the layered architecture, the layers can communicate only with the adjacent layers. Teams needing structured development.



12.3.1 Layered architecture - variations

12.3.2 Adding layers

New layers to enforce architectural constraints, helps control access and dependencies. Example: add *service layer* for share common logic

12.3.3 Isolation

Each layer is independent, interacting only through defined contracts. Prevents tight coupling and cross-layer dependencies. The risk is the **Sinkhole antipattern**.

12.3.4 Open layers

Allow direct access to any layer for efficiency, useful for simple operations. Improves performance but reduces architectural control.

12.3.5 Open vs close layers

- Closed layers → stability, modularity, easier changes
- Open layers → flexibility, performance, but higher risk of coupling

Synkholes acceptable threshold: 20% manageable, 80% change architecture style!

	Characteristic	Rating	Notes
Structural	Simplicity	Very High	
	Modularity	Very low	
Engineering	Maintainability	Very low	Changes can impact every layer
	Testability	Low	Mock/stub components
	Deployability	Very low	At every change !
	Evolvability	Very low	Changes can impact every layer
Operational	Responsiveness	Medium	Caching and multithreading
	Scalability	Very low	Lack of modularity
	Elasticity	Very low	Lack of modularity
	Fault tolerance	Very low	Faults impact the whole system, all layers must be restarted

When to use	When not to use
Applications with clear separation	Not ideal for high-performance or highly dynamic systems
Small, simple applications	Large and complex applications
As starting point for situations with low budget and short time	
To quickly test new ideas	
Examples of usage: operating systems, desktop applications	

12.4 Modular monolithic style

Domain partitioned: each module encapsulates a business capability. Strong internal boundaries required. Modules with clear responsibilities. Deployed as a single unit of software. E.g. Web archive (WAR) file.⁷

Modules might have their own dedicated repositories.


In this we decoupled all the modules like Customer, Order, etc. The team is also organized in different way before a team for each layer, now we have a team for each modules or a pool of modules.

Data topology

Also in this style we can have single database or more domain-specific databases for the best performance.

Interaction between modules


The modules **communicate directly** (peer to peer) or **through a mediator** the goal is orchestrate different modules.



12.4.1 Layered vs modular: components organization

- Layered: `com.app.presentation.customer.profile`
- Modular: `com.app.customer.profile`; `com.app.customer.profile.presentation`; ...

12.4.2 Variations of modular monolithic style




12.4.3 Advantages and disadvantages of modular monolithic style

Key characteristics	Tradeoffs:
Simpler operational model than microservices	Risk of becoming tightly coupled over time
No network overhead	Limited independent scalability
Easier testing and debugging with respect to layered style	Large codebase can slow development
Requires discipline to maintain boundaries	Harder team autonomy at scale
	Risk of long start up

12.5 Pipeline architecture


Also known as Pipes and Filters. Technically partitioned. Data flows through a sequence of processing steps. Each step transforms input into output. Filters are independent and stateless (ideally)



	Characteristic	Modular Monolith	Notes
Structural	Simplicity	Very High	
	Modularity	Low	
Engineering	Maintainability	Low	Better than Layered because of higher modularity
	Testability	Low	As above
	Deployability	Low	As above
	Evolvability	Low	As above
Operational	Responsiveness	Medium	
	Scalability	Very low	
	Elasticity	Very low	
	Fault tolerance	Very low	
When to use		When not to use	
Early-stage systems needing structure		Systems with high demands for operational characteristics	
Situations with tight time and budget constraints		Systems requiring independent scaling	
Teams wanting simplicity over distribution		When majority of changes are technically-oriented	
Teams working with DDD			
Applications with strong domain boundaries			

12.5.1 Components


- **Filters:** processing units performing specific business functions
- **Pipes:** connectors passing data between filters, unidirectionally. Can be deployed as remote service (style remains monolith). Can be either synchronous or asynchronous (e.g., threads, messages).
- **Data format:** must be consistent or transformed
- **Database:** usually a single db, but filter can have their own database too
- **Producer:** starting point of a process
- **Transformer:** enhance, transform, compute (map)
- **Tester:** test input w.r.t. to criteria (reduce), optionally produces output
- **Consumer:** termination point, with persistence



Key characteristics:	Tradeoffs:
Easy to extend by adding filters	Latency due to multiple stages
Clear separation of concerns	Harder debugging across pipeline
Examples: Bash, Zsh, MapReduce	Weak fit for complex business logic
	Data format coupling between filters


12.6 Micro-kernel - plug-in architecture style

Also known as Plug-in architecture. Core system with basic and stable functionalities («happy paths»). It can be technically or domain partitioned. Plug-ins extend system capabilities and incorporate complexity. Registry or discovery mechanism. Data topology: one db or plug-in specific data sources



	Characteristic	Pipeline	Notes
Structural	Simplicity	High	Separation of concerns in filters
	Modularity	Low	
Engineering	Maintainability	Low	Separation of concerns in filters
	Testability	Medium	
	Deployability	Low	
	Evolvability	Medium	
Operational	Responsiveness	Medium	Separate deployments can improve it, but at the expense of simplicity and cost.
	Scalability	Very low	
	Elasticity	Very low	As above
	Fault tolerance	Very low	As above
When to use		When not to use	
Data processing and transformation pipelines		In systems with non-deterministic flows (use event-driven a.s.)	
Streaming or ETL systems		When fault tolerance is important	
If properly implemented: independent, reusable processing steps		When bi-directional communication is needed	
		If high demands of operational characteristics	

12.6.1 UI: variants



12.6.2 Plug-in components

- They should be independent and self-contained
- Communication p2p with the core system, not between plugins
- Compile based vs runtime base
- Common semantic, e.g. `app.plugin.<domain>.<context>` -> Example. `app.plugin.editors.plain`
- Plug-in could a standalone remote service
 - The architecture remains monolithic due to coupling with core system
 - Advantages: scalability, throughput, asynchronous communication
 - Disdvantages: reliability, costs, complexity

12.7 Monolithic Styles - recap

Characteristics:	Trade-offs:
Highly extensible and customizable	Performance overhead from indirection
Separation between core and features	Complex plug-in lifecycle management
Supports product variability	Testing combinations of plugins can be hard
Stable core is critical	Core design mistakes are costly

When to use	When not to use
Common Product-based systems	Core systems that need frequent changes
Systems requiring extensibility	High dependencies between plugins
Platforms with plugin ecosystems	
Customization-driven applications	
Examples: browsers, IDEs, PMD, Jira	

	Characteristic	Microkernel	Notes
Structural	Simplicity	High	
	Modularity	Medium	Extensions easier with plugins
Engineering	Maintainability	Medium	Functionalities are better isolated w.r.t other monolithic styles
	Testability	Medium	As above
	Deployability	Medium	As above
	Evolvability	Medium	Extensions easier with plugins
Operational	Responsiveness	Medium	Microkernel applications are generally small
	Scalability	Very low	
	Elasticity	Very low	
	Fault tolerance	Very low	

	Characteristic	Layered (n-tier)	Modular Monolith	Pipeline	Microkernel
Structural	Simplicity	Very High	Very High	High	High
	Modularity	Very low	Low	Low	Medium
Engineering	Maintainability	Very low	Low	Low	Medium
	Testability	Low	Low	Medium	Medium
	Deployability	Very low	Low	Low	Medium
	Evolvability	Very low	Low	Medium	Medium
Operational	Responsiveness	Medium	Medium	Medium	Medium
	Scalability	Very low	Very low	Very low	Very low
	Elasticity	Very low	Very low	Very low	Very low
	Fault tolerance	Very low	Very low	Very low	Very low

Chapter 13

Software Quality


ISO/IEC 9126: issued in 1991 and revised in 2001, being retired. Iso/IEC 25XX - SQuaRE: Software Product Quality Requirements and Evaluation. Is a family of standards in development.

Software quality is divided in two parts: Product quality (internal, external, in use, data) and process quality.



13.1 Target entities

Product quality addresses several different entities in a complex information system



13.2 Model structure

- **Characteristics:** main aspects eg. usability
- **Sub-characteristics:** specific aspects eg. accessibility
- **Metrics:** evaluate a specific (sub)-characteristic, eg. number of clicks to reach a page
- **Measure elements:** Fundamental

13.3 Software product quality

- Functional suitability
- Performance efficiency
- Compatibility
- Interaction capability (Usability)
- Reliability
- Security
- Maintainability
- Flexibility

13.3.1 Functional suitability

Capability of a product to provide functions that meet stated and implied needs of intended users when it is used under specified conditions: completeness, correctness, appropriateness.

13.3.2 Performance efficiency

Capability of a product to perform its functions within specified parameters and be efficient in the use of resources under stated condition: time behaviour, resource utilization, capacity.

13.3.3 Compatibility

Capability of a product or its component to exchange information with other products, systems or components, and/ or perform its required functions, while sharing the same hardware or software environment: co-existence, interoperability.

13.3.4 Interaction capability (Usability)

Capability of a product to be used by specified users to achieve specified goals with effectiveness, efficiency and satisfaction in a specified context of use: Recognizability, learnability, operability.

13.3.5 Reliability

Capability a product, to perform specified functions under specified conditions for a specified period of time: maturity, availability, fault tolerance, recoverability.

13.3.6 Security

Capability of a product to protect information and data so that persons or other products or systems have the degree of data access appropriate to their types and levels of authorization, and to defend against attack patterns by malicious actors: confidentiality, integrity, integrity.


13.3.7 Maintainability

Capability of a product to be modified by the intended maintainers with effectiveness and efficiency: modularity, reusability, analysability, modifiability, testability.

13.3.8 Flexibility

Capability of a product to serve a different or expanded set of requirements or contexts of use; or the ease with which the product can adapt itself to changes in its requirements, contexts of use, or system environment: adaptability, installability, replaceability.

13.3.9 Sub-characteristics




13.3.10 Usability - User error protection

ID	UEp-1-G
Name	Avoidance of user operation error
Description	What portion of user actions and inputs are protected against causing any system malfunction?
Measurement function	$X = A/B$, where $A = \#$ user actions and inputs protected against causing any system malfunction; $B = \#$ user actions and inputs that could be protected
Interpretation	$0 \leq X \leq 1$ The closer to 1 the better

13.3.11 Reliability - Maturity

ID	RMa-1-G
Name	Test coverage
Description	How much of the required test cases has been executed during testing?
Measurement function	$X = A/B$, where $A = \#$ test cases actually performed representing operation scenario during testing; $B = \#$ test cases to be performed to cover requirements
Interpretation	$0 \leq X \leq 1$ The closer to 1 the better

ID	MMo-2-S
Name	Cyclomatic Complexity
Description	How many software modules have the acceptable cyclomatic complexity?
Measurement function	$X = A/B$, where $A = \#$ of software modules which have the cyclomatic complexity exceeds a threshold; $B = \#$ of software modules in a product. Threshold is defined by each project or organization and is possibly different value by a programming language, or type of module
Interpretation	$0 \leq X \leq 1$ The closer to 1 the better



13.3.12 Maintainability - Modularity

13.4 Quality in Use

13.4.1 Quality in Use - Characteristics

- Effectiveness: accuracy and completeness with which identified goals are achieved
- Efficiency: resources used in relation to the results achieved
- Satisfaction: extent to which a user or other stakeholder perceives the results of use of a system, product or service to meet their needs and expectations
- Freedom from risk: degree to which the quality of a product or system mitigates or avoids potential risk to the user, organisation or project, including risks to economic status, human life, health, or the environment.
- Context coverage: degree to which a product or system can be used with effectiveness, efficiency, satisfaction, and freedom from risk in both specified contexts of use and in contexts beyond those initially explicitly identified.

13.4.2 Maintainability - Effectiveness

ID	MMo-2-S
Purpose	Cyclomatic Complexity
Method of application	How many software modules have the acceptable cyclomatic complexity?
Definition	$M1 = 1 - \sum A_i $, where $A_i = \#$ Proportional value of each missing or incorrect component in the task output
Interpretation	$0 \leq M1 \leq 1$ The closer to 1 the better

13.4.3 Quality in Use Influences

- Influence by use: usability and accessibility
- Influence by use of output: Perceived usefulness and Avoidance of harm
- Influence of context: Consequence of use
- Context coverage: Context completeness and Context extensibility

13.5 Data quality

Characteristics from two perspectives:

- **Inherent:** Focus on the data itself, their structure, schema, and semantics
- **System dependent:** Focus on how the data is stored and processed

Inherent	Inher. / Sys. Dep.	Sys. Dep.
Accuracy	Accessibility	Availability
Completeness	Compliance	Portability
Consistency	Confidentiality	Recoverability
Currency	Efficiency	
Credibility	Precision	
	Understandability	
	Traceability	

13.5.1 Accuracy

Correspondence between data and reality

- **Syntactic:** Value belongs to a set of validated information
- **Semantic:** The meaning, the content, corresponds to the reality

Accuracy: Open vs. Closed world

- Closed World Assumption (CWA): The knowledge represented in the data (and its schema) is complete. E.g., if a code appears in the list of valid codes it is accurate, otherwise it is wrong.
- Open World Assumption (OWA): The knowledge about the data is (knowingly) incomplete E.g., if a code is in the list of valid ones it is accurate, otherwise it is not possible to immediately decide

Chapter 14

Software Observability

Observability: is the ability to understand the internal state of a system by examining its outputs.

Telemetry data - i.e. logs and metrics allows:

- detecting possible problems
- troubleshooting and handle novel problems (unknown unknowns)
- understanding the motivation for system behavior (why is this happening?)

14.1 OpenTelemetry

OpenTelemetry is an observability framework and toolkit designed to facilitate the generation, export and collection of telemetry data such as traces, metrics, and logs. Open source, as well as vendor- and tool-agnostic.

14.1.1 Key concepts

- **Telemetry:** data emitted by a system during operation, requires instrumentation of the system
- **Service Level Indicator (SLI):** measures the behavior of the system and captures the user perspective
- **Service Level Objective (SLO):** attaches business value to one or more SLIs

14.1.2 Distributed Tracing

Distributed tracing observe requests as they propagate through complex, distributed systems. It is based on three main types of signals: **Logs**, **Spans** and **Traces**.

Logs

A log is a timestamped message emitted by services or other components. Built-in logging facilities in most languages, often lack contextual information.

Log format

- Structured: consistent schema or typed fields, can be JSON or other
- Unstructured: don't follow a consistent structure, therefore much more difficult to parse and analyze at scale
- Semistructured: include machine-readable key/value pairs or delimited fields but do not guarantee a stable schema

Common Log Format (CLF) 192.168.1.30 - alice [01/May/2025:07:20:10 +0000] "GET /index.html HTTP/1.1" 200 9481

14.1.3 Span

A span represents a single unit of work or operation. Spans describes specific individual operations triggered by a (user) request. It contains:

- name
- time-related data
- structured log messages
- other metadata (i.e. Attributes)

Span Kind

- Client: outgoing remote call such as an outgoing HTTP request or database call
- Server: incoming remote call such as an incoming HTTP request or remote procedure call
- Internal: operations which do not cross a process boundary
- Producer: creation of a job which may be asynchronously processed later
- Consumer: processing of a job created by a producer and may start long after the producer span has already ended


14.1.4 Trace

A (distributed) trace records the path taken by a single request as it propagates through multiple services/components.

A trace is made of one or more spans, the first span represents the *root span*. (span represents a request from start to finish). The spans underneath the parent provide a more in-depth context of what occurs.

14.2 Context Propagation

All spans part of the same trace must be identified and put in relation. Context information is passed to called components Standard traceparent HTTP header field: version, trace-id, parent-id, trace-flags



14.3 Metrics

A metric is a measurement of a service captured at runtime. Metric event: measure values, timestamps and other metadata (i.e. Attributes).

Capture a measurement: Name, Kind, Unit, Description.

14.3.1 Metric kind

- Counter: A value that accumulates over time
- UpDownCounter: A value that accumulates over time, but can also go down again, e.g. queue length
- Gauge: Measures a current value at the time it is read
- Histogram: A client-side aggregation of values, such as request latencies

Asynchronous: used if you access available only to the aggregated/current value


Entities Example App, browser, container, host, process, thread, service, etc.

Metrics examples For entity process:

Name	Type	Unit	Description
cpu time	Counter	s	Total CPU seconds (by mode)
cpu utilization	Gauge	1	$\frac{\Delta time}{(elapsed * \#CPU)}$
memory usage	UpDownCounter	by	Amount of physical memory in use

For entity httpserver:

Name	Type	Unit	Description
request duration	Histogram	s	Duration of requests (attributes: method, url, ...)
active_requests	UpDownCounter	request	Number of active requests (attributes: method, url, ...)
request body size	Histogram	by	Payload size in bytes (attributes: method, url, ...)



Metric Example

- Resource: namespace: my-site, name: frontend-web, id:http-server-136
- Metric: active_requests
- Data point: value: 35, method: GET, url: /index.html

Model

14.4 Profiling

A running program consumes computational resources such as memory, CPU cycles, and elapsed time. To **profile** a program is to measure the consumption of such resources by specific elements of the program.

14.4.1 Why profiling?

Profiling means:

- make program more efficient
- developers more productive
- by identifying which program elements to optimize

The risk without profiling is to waste time in optimizing a method consuming few resources resulting in little impact.

14.4.2 Sampling vs. Tracing

Two main approaches to collect profiling data: **Sampling**, an external harness samples the status of the program and **Tracing**, an instrumented program emits samples during execution.

Sampling

Execution sampling provides an approximation of CPU profiling.

- Periodically interrupts the program (e.g., every 1ms)
- Records what the CPU is doing at that moment
- Low overhead, Works in production
- Not perfectly precise, but statistically accurate

Tracing

Tracing - a.k.a. instrumentation profiling - records function calls (all or selected ones).

- High detail (exact timings, call counts)
- Higher overhead
- Can distort performance

14.4.3 Profile Types

- On-CPU
- Stall cycles: block on processor or hardware resources — typically memory I/O
- Memory: events related to memory growth e.g. malloc(), mmap(), etc.
- I/O: issuing of I/O, such as file system, storage device, and network
- Off-CPU: time spent while a thread is blocked, not running on a CPU (off-CPU time), e.g., waiting on I/O, locks, timers, a turn on-CPU, waiting for paging or swapping.

14.4.4 Flamegraph


Collection of stack samples represented as a column of boxes, each box represents a function (a stack frame).

The y-axis shows the stack depth, from root at the bottom to leaf at the top. The top box is the function that was on-CPU

The x-axis spans the stack trace collection (not time passing). Ordering is alphabetical on the function names. When identical boxes are adjacent, they are merged.

The width of box is the frequency at which that function was present in samples. Functions with wide boxes were more present in the stack traces than those with narrow boxes, in proportion to their widths.

The background color is random to help the eye differentiate boxes.



captionoffigureExample of flamegraph

Flamegraph interpretation

The top is the function that was on-CPU when the stack trace was collected. Look for large plateaus along the top edge: this means a single function was frequently running on-CPU. Reading top down shows ancestry. Reading bottom up shows code flow and the bigger picture

The width of function boxes can be directly compared, if a function box is wider than another, this may be because it consumes more CPU per function call or that the function was simply called more often.

Major forks in the flame graph can indicate, a logical grouping of code (work in stages), a conditional statement.

14.5 Java Flight Recorder (JFR)

JFR is a low-overhead profiling and monitoring system built into the JVM. It records events about CPU, memory, GC, threads, I/O, and JVM internals. Data is stored in a compact .jfr file for later analysis. It supports both: continuous production monitoring and short profiling sessions.

14.5.1 Collected Events

- JVM Internals (e.g. CPU, Threads, Monitors)
- Profiling (execution sampling)
- Garbage Collection
- Memory (object allocation)
- Exception
- I/O
- Class loading

14.5.2 JFT settings

Two predefined settings:

- default: No jdk.ExecutionSample; Focus on: GC, locks, I/O, JVM internals; Very low overhead
- profile: Enables: jdk.ExecutionSample + allocation profiling; Higher overhead (1-2%); Suitable for flamegraphs

JFR profiling session example An app is run with JFR started from begin:

<pre> Event Type Count Size (bytes) ===== jdk.NativeLibrary 1516 129323 jdk.ModuleExport 1017 10447 jdk.SystemProcess 1002 104888 jdk.BooleanFlag 966 28708 jdk.ActiveSetting 752 18683 jdk.GCPhaseParallel 319 7794 jdk.ModuleRequire 312 2964 </pre>	<pre> java \ -XX:StartFlightRecording=filename=profiling-recording.jfr -cp target/classes \ it.polito.swda.profiling.Example jfr summary profiling-recording.jfr </pre>
--	--

Chapter 15

Distributed Architectural Styles


15.0.1 Service-based Architecture

Distributed macro-layered structure: UI+ services

Quanta: UI separately deployed. Separately deployed coarse-grained services.

Optionally, monolithic database

Domain services are independent, usually do not communicate each other



15.0.2 Domain services design



Figure 15.1: Technical partitioning and Domain partitioning

15.0.3 API access, UI and DB options

API access façade acts as orchestrator: failures can be handled at this level.

UI/DB options: Single monolithic, Domain-based or Service-based.

API gateway options between UI and domain services.




Figure 15.2: UI options




Figure 15.3: DB options


15.0.4 Database logical partitioning

Using a single shared library for DB entity objects impacts all services in cases of changes. Logical partitioning and multiple share libraries for db entity objects reduces coupling. It should be as fine-grained as possible

Using a single shared library for DB entity objects impacts all services in cases of changes. Logical partitioning and multiple share libraries for db entity objects reduces coupling. It should be as fine-grained as possible.

When to Use

- Small-medium systems
- Limited scalability needs
- No or low need of service coordination
- Teams new to distributed systems
- As intermediate steps towards migration to other distributed styles



15.0.5 Service-based architecture: characteristics

	Characteristic	Rating	Notes
Structural	Simplicity	Medium	
	Modularity	Medium	Domain scoping, federated UI and BD possible
Engineering	Maintainability	High	As consequence of modularity
	Testability	High	As above
	Deployability	High	As above
	Evolvability	High	As above
Operational	Responsiveness	Medium	
	Scalability	Medium	Coarse-grained services
	Elasticity	Low	
	Fault tolerance	Medium	Self-contained services

15.1 Microservices architecture

Domain partitioning took to the extreme. Heavily inspired by the idea of domain-driven-design and DevOps
Fundamental concept of “bounded context”

- Everything related to a portion of the domain is visible internally but opaque to other bounded contexts
- It drives systematically interventions of decoupling

Protocol-aware heterogenous interoperability

- Standardized communications (e.g., REST)
- Independent services
- Inter-service communication possible




Figure 15.4: Overview




Figure 15.5: Monitoring and governance: sidecar pattern

15.1.1 Granularity of services

Extreme granularity may affect performances

Main decision criteria for micro-services:

- Purpose (i.e., problem domain)
- (Avoid/minimize) transactions
- Choreography

Other aspects:

- Data isolation: per service or per domain
- API layer “neutrality”
- Frontends: monolithic vs coupling with backend services
- Choreography vs orchestration

	Characteristic	Rating	Notes
Structural	Simplicity	Very low	
	Modularity	Very high	
Engineering	Maintainability	Very high	
	Testability	Very high	
	Deployability	Very high	
	Evolvability	Very high	
Operational	Responsiveness	Low	Many network calls to complete work; data caching and replication + choreography can improve it
	Scalability	Very high	
	Elasticity	High	
	Fault tolerance	Very high	

15.1.2 Microservices architecture: characteristics

15.2 Event-driven architecture

Main implementation: choreographed and technically partitioned Decoupled event processing components asynchronously trigger and responds via events. Asynchronous fire-and-forget communication. Loose coupling.

Four primary components:

- Initiating event
- Event broker
- Event processor
- Derived events




Figure 15.6: Event-driven architecture





Figure 15.7: Fire and forget



EDA processing - example with SPG

15.2.1 Events vs messages

15.2.2 Further options

- **Triggering extensible events:** It increases evolvability of the systems
- **Asynchronous vs synchronous = performance vs responsiveness:** NB asynchronous communication decouples architectural quanta
- Broadcast capabilities for implementing semantic decoupling
- Granularity of events: Be aware of swarm of gnats antipattern !
- Error handling through Workflow Event pattern or tools

	EVENTS	MESSAGES
ACTION	Reacting to something already happened	Obeying to a request with something that needs to be done
ANSWER	It doesn't require an answer	It -typically- requires an answer
COMMUNICATION STYLE	One-to-many (e.g., publish and subscribe)	One-to-one (or point-to-point)
COMMUNICATION CHANNEL	Topic, stream, notification service	Queue, messaging service

15.2.3 Event payload

Data-based payload

- It avoids pression on DB: payload contains all information for subscribers
- Extra support for data consistency and integrity
- Data format creates coupling: Stamp coupling
- Versioning: Data and Contract

Key-based payload

- Payload contains pairs of keyvalues to later query a DB. Typically, JSON or XML, Human readable
- DB remains the single source of information
- Minimum contracts and network usage
- No need of versioning




Figure 15.8: Data-based payload, Key-based payload


Event payload - Example

15.2.4 Events payload: tradeoffs

	DATA-BASED	KEY-BASED
Performance and scalability	😊	😞
Contract management	😞	😊
Stamp coupling	😞	😊
Bandwidth utilization	😞	😊
Restricted db access	😊	😞
Overall system fragility	😞	😊

15.2.5 Data loss prevention

1. While A publishes an event, it crashes before ack from event broker
2. B accepts an event, but it crashes before processing it
3. A data error prevents B from ensuring persistence



15.2.6 Request-reply processing

Also known as *pseudosynchronous communications*. It requires a request queue and a reply queue. Event producers wait until confirmation of event producer is retrieved from the reply queue.

Two common implementations:

- Correlation ID
- Temporary queue with specific id

15.3 Orchestrated EDA


Orchestration vs Choreography.

Mediation by a specific component/service.

Event mediator typically uses messages. The mediator is responsible of catching events and generate specific point-to-point derived events. Usually with queues.

Without further communicating its own actions (fire and forget principle)

Orchestrator's options: monolithic, domain-based, service-dedicated.



15.3.1 Event drive architecture: usage

When to use

- Complex workflows, with thousands scenarios
- Dynamic user content
- Reactive systems
- Real time processing

When NOT to use

- High needs of synchronous communication
- Static coupling (contracts) increases
- Dynamic coupling (sync calls) increases

15.3.2 EDA architecture: characteristics


	Characteristic	Rating	Notes
Structural	Simplicity	Low	Dependent also on dynamic coupling and data topologies
	Modularity	High	
Engineering	Maintainability	High	Due to non-deterministic workflow
	Testability	Low	
	Deployability	Medium	
	Evolvability	Very high	
Operational	Responsiveness	Very high	Additional processors can be added; db still bottleneck
	Scalability	High	
	Elasticity	Medium	
	Fault tolerance	Very high	

15.4 Space-based architecture

Technically partitioned. Complex implementation. Designed to optimize operations with large volumes of data, users, and load peaks.

Main features:

- Eliminates database bottleneck
- Distributed data management
- Distributed processing units



15.4.1 Space-based architecture: elements

15.4.2 Replicated cache data: collisions

Collisions typically occur when the update rate of the data in the cache exceeds the replication latency.


Several factors influence the amount of data collisions:

Element name	Description
Processing units	Contain the application functionality
Virtualized middleware	The collection of infrastructure-related artifacts that are used to manage and coordinate the processing units
Messaging grid	It manages input requests and session states (e.g., a web server with load balancing)
Data grid	Manages the synchronization and replication of data between processing units. It can be implemented also in processing units.
Processing grid	Used to manage request orchestration between multiple processing units; optional
Deployment manager	Manages the starting and tearing down of processing unit instances based on load
Data pumps	Asynchronously send updated data to the database. It decouples processing units and database. Usually implemented with messaging and queues. Can be mapped to domains
Data writers	Perform the updates from the data pumps. They can be multiple and dedicated to data pumps
Data readers	Read database data and deliver it to processing units (upon startup) through reverse data pumps.

- N : # of processing unit instances containing the same cachce
- UR: cache update rate (updates/seconds)
- S: cache size (# of rows)
- RL: Replication latency (milliseconds)

$$\text{Collision Rate} = N \times \frac{UR^2}{S} \times RL$$

15.4.3 Distributed caching



15.4.4 Replicated vs distributed caching

= Asynchronous replication vs synchronous replication Replicated caching:

- requires in-memory data grid, very fast
- difficult with very high data volumes and frequent updates
- possible also with near-cache techniques: a caching server + partial front caches

Criteria	Replicated	Distributed
Optimization	Performance	Consistency
Cache size	Small	Large
Type of data	Mostly static	Highly dynamic
Update frequency	Low	High
Fault tolerance	High	Low

15.5 SBA architecture - Usage

15.5.1 When to use:


- High volume applications
- Large, concurrent user load
- Low-latency systems

15.5.2 Space based architecture: characteristics

	Characteristic	Rating	Notes
Structural	Simplicity	Very low	
	Modularity	Medium	Technically partitioned
Engineering	Maintainability	Medium	
	Testability	Very low	Difficult to simulate very high loads
	Deployability	Medium	
	Evolvability	Medium	
Operational	Responsiveness	Very high	
	Scalability	Very high	
	Elasticity	Very high	
	Fault tolerance	Low	Technically partitioned, distributed computing and data management


15.6 Orchestration-driven service-oriented architecture

- Enterprise-level architecture
- Layers with specific responsibilities
- Strong governance through taxonomy of services
- Services share enterprise bus
- Usage in companies decreasing, due to extreme technical partitioning



15.6.1 Service-oriented architecture: layers

- **Business services:** no business logic, no code, only I/O schemas (service signatures)
- **Enterprise service bus:** Transaction coordination and message transformation. Also for internal calls, it can “eat” the whole architecture
- **Enterprise services:** Building blocks of isolated and combinable/reusable functionalities
- **Application services:** one-off, single implementation services, typically not reusable
- **Infrastructure services:** Monitoring, authentication, logging



Reasons for decay after the 2010s Domains concepts spread so thinly that they were virtually ground to dust

15.6.2 ODSA architecture: characteristics

	Characteristic	Rating	Notes
Structural	Simplicity	Very low	
	Modularity	High	
Engineering	Maintainability	Very low	Because of very high coupling
	Testability	Very low	As above
	Deployability	Very low	
	Evolvability	Very low	Ripple effects for any functional/domain change
Operational	Responsiveness	LOW	
	Scalability	High	Through session duplication
	Elasticity	Medium	
	Fault tolerance	Medium	

15.7 Distributed Styles Architectural Characteristics: synopsis

15.7.1 Structural and engineering characteristics

	Characteristic	Service based	Microservice	Event driven	Space based	Orch. Driven Service or.
Structural	Simplicity	Very low	Very low	Low	Very low	Very low
	Modularity	Very high	Very high	High	Medium	High
Engineering	Maintainability	Very high	Very high	High	Medium	Very low
	Testability	Very high	Very high	Low	Very low	Very low
	Deployability	Very high	Very high	Medium	Medium	Very low
	Evolvability	Very high	Very high	Very high	Medium	Very low

15.7.2 Operational characteristics

	Characteristic	Service based	Microservice	Event driven	Space based	Orch. Driven Service or.
Operational	Responsiveness	Medium	Low	Very high	Very high	Low
	Scalability	Medium	Very high	High	Very high	High
	Elasticity	Low	High	Medium	Very high	Medium
	Fault tolerance	Medium	Very high	Very high	Low	Medium

Chapter 16

Deployment

16.1 Deployment strategies

16.1.1 Infrastructure deployment

- On premise: Software system installed and run on servers physically located in the organization
- Hosted: Software system installed on third parties centers
- Cloud: Software system is virtualized over the cloud

16.1.2 Application deployment strategies

Goal: define how an application, or new versions of it, are deployed on an infrastructure

Main strategies: In place upgrade, Blue-Green, Canary / Rolling

In place upgrade


Installation on top of current version. It requires downtime. All existing settings are preserved

Blue-green

Two environments, a load balancer directs traffic. Blue env.: where the current application runs. Green: idle

When a new version is ready for deployment:

- Update deployment, test and monitoring on the green environment
- When situation is deemed stable, all traffic is sent to the green env.
- The green environment now is the main one, the blue is idle




16.1.3 Canary

"Canary in the coal mine".

Initially, only a few nodes hosts the new deployed application version. Load balancer forwards a minority fraction of requests to new version. Performances are monitored. If there are no severe issues:

- load balancer increases the percentage towards new application version;
- The new application is deployed in further nodes

This incremental process continues until the new version is handling all traffic on the majority of nodes, and few ones remains idle, so that process can restart with further new versions. In case of problems, balance can be temporary reversed



Canary with feature flags

- New features are flagged in the code, so that they can be enabled via configuration files
- New features will be enabled for a small groups of users
- Monitor the performances of the “canary” group
- If no significant issues arise, increase dimension of “canary” group
- Progressively all users will use the new features
- In case of problems, dimensions of canary group can be temporary diminished

Rolling with feature flags

- Same as canary, but exposure of new changes is made on nodes, not users
- New features are flagged in the code, so that they can be enabled via configuration files
- New features will be enabled for a small groups of (users) nodes
- Monitor the performances of the “canary” group
- If no significant issues arise, increase dimension of “canary” group
- Progressively all (users) nodes will use the new features
- In case of problems, dimensions of canary group can be temporary diminished

16.2 Deployment visualization

Deployment diagram

Goal: describe the mapping of software systems and containers (and any other deployable unit) onto deployment nodes

Method: C4 model with further boundaries for each deployment node

Deployment node: a device or execution environment

- Physical infrastructure (e.g. a physical server or device)
- Virtualised infrastructure (e.g. a virtual machine, IaaS)
- Containerised infrastructure (e.g. a Docker container)
- An execution environment (e.g. a database server)



Temporary page!

L^AT_EX was unable to guess the total number of pages correctly. As there was some unprocessed data that should have been added to the final page this extra page has been added to receive it.

If you rerun the document (without altering it) this surplus page will go away, because L^AT_EX now knows how many pages to expect for this document.